

# Apache Kafka vs. Traditional Middleware (MQ, ETL, ESB)

Friends, Enemies or Frenemies?



Kai Waehner

Technology Evangelist

kontakt@kai-waehner.de

LinkedIn

@KaiWaehner

[www.confluent.io](http://www.confluent.io)

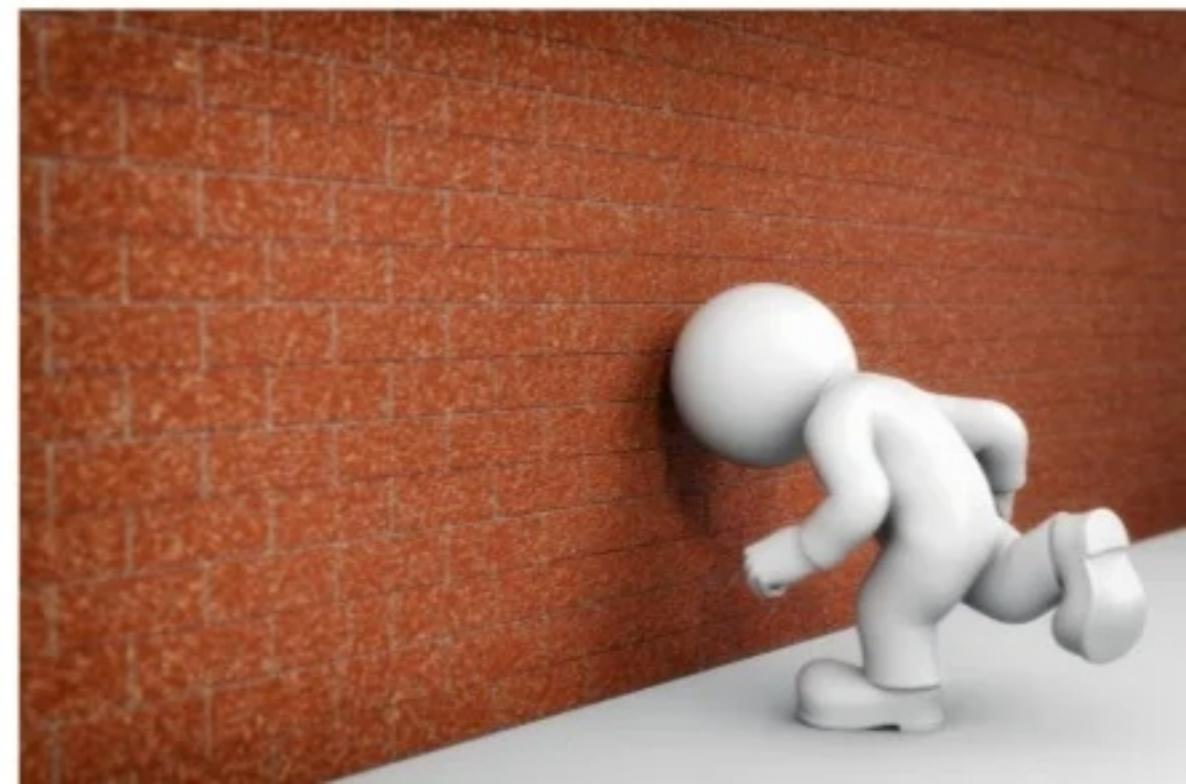
[www.kai-waehner.de](http://www.kai-waehner.de)



## Agenda

---

1. Traditional Middleware
2. Event Streaming Platform
3. Enemies
4. Friends
5. Frenemies



## Agreement:

Kafka is the **de facto standard** for ...

- messaging at scale!
- decoupling of microservices!
- reliable, lightweight stream processing!

## Controversial discussion:



## 1. Traditional Middleware

2. Event Streaming Platform

3. Enemies

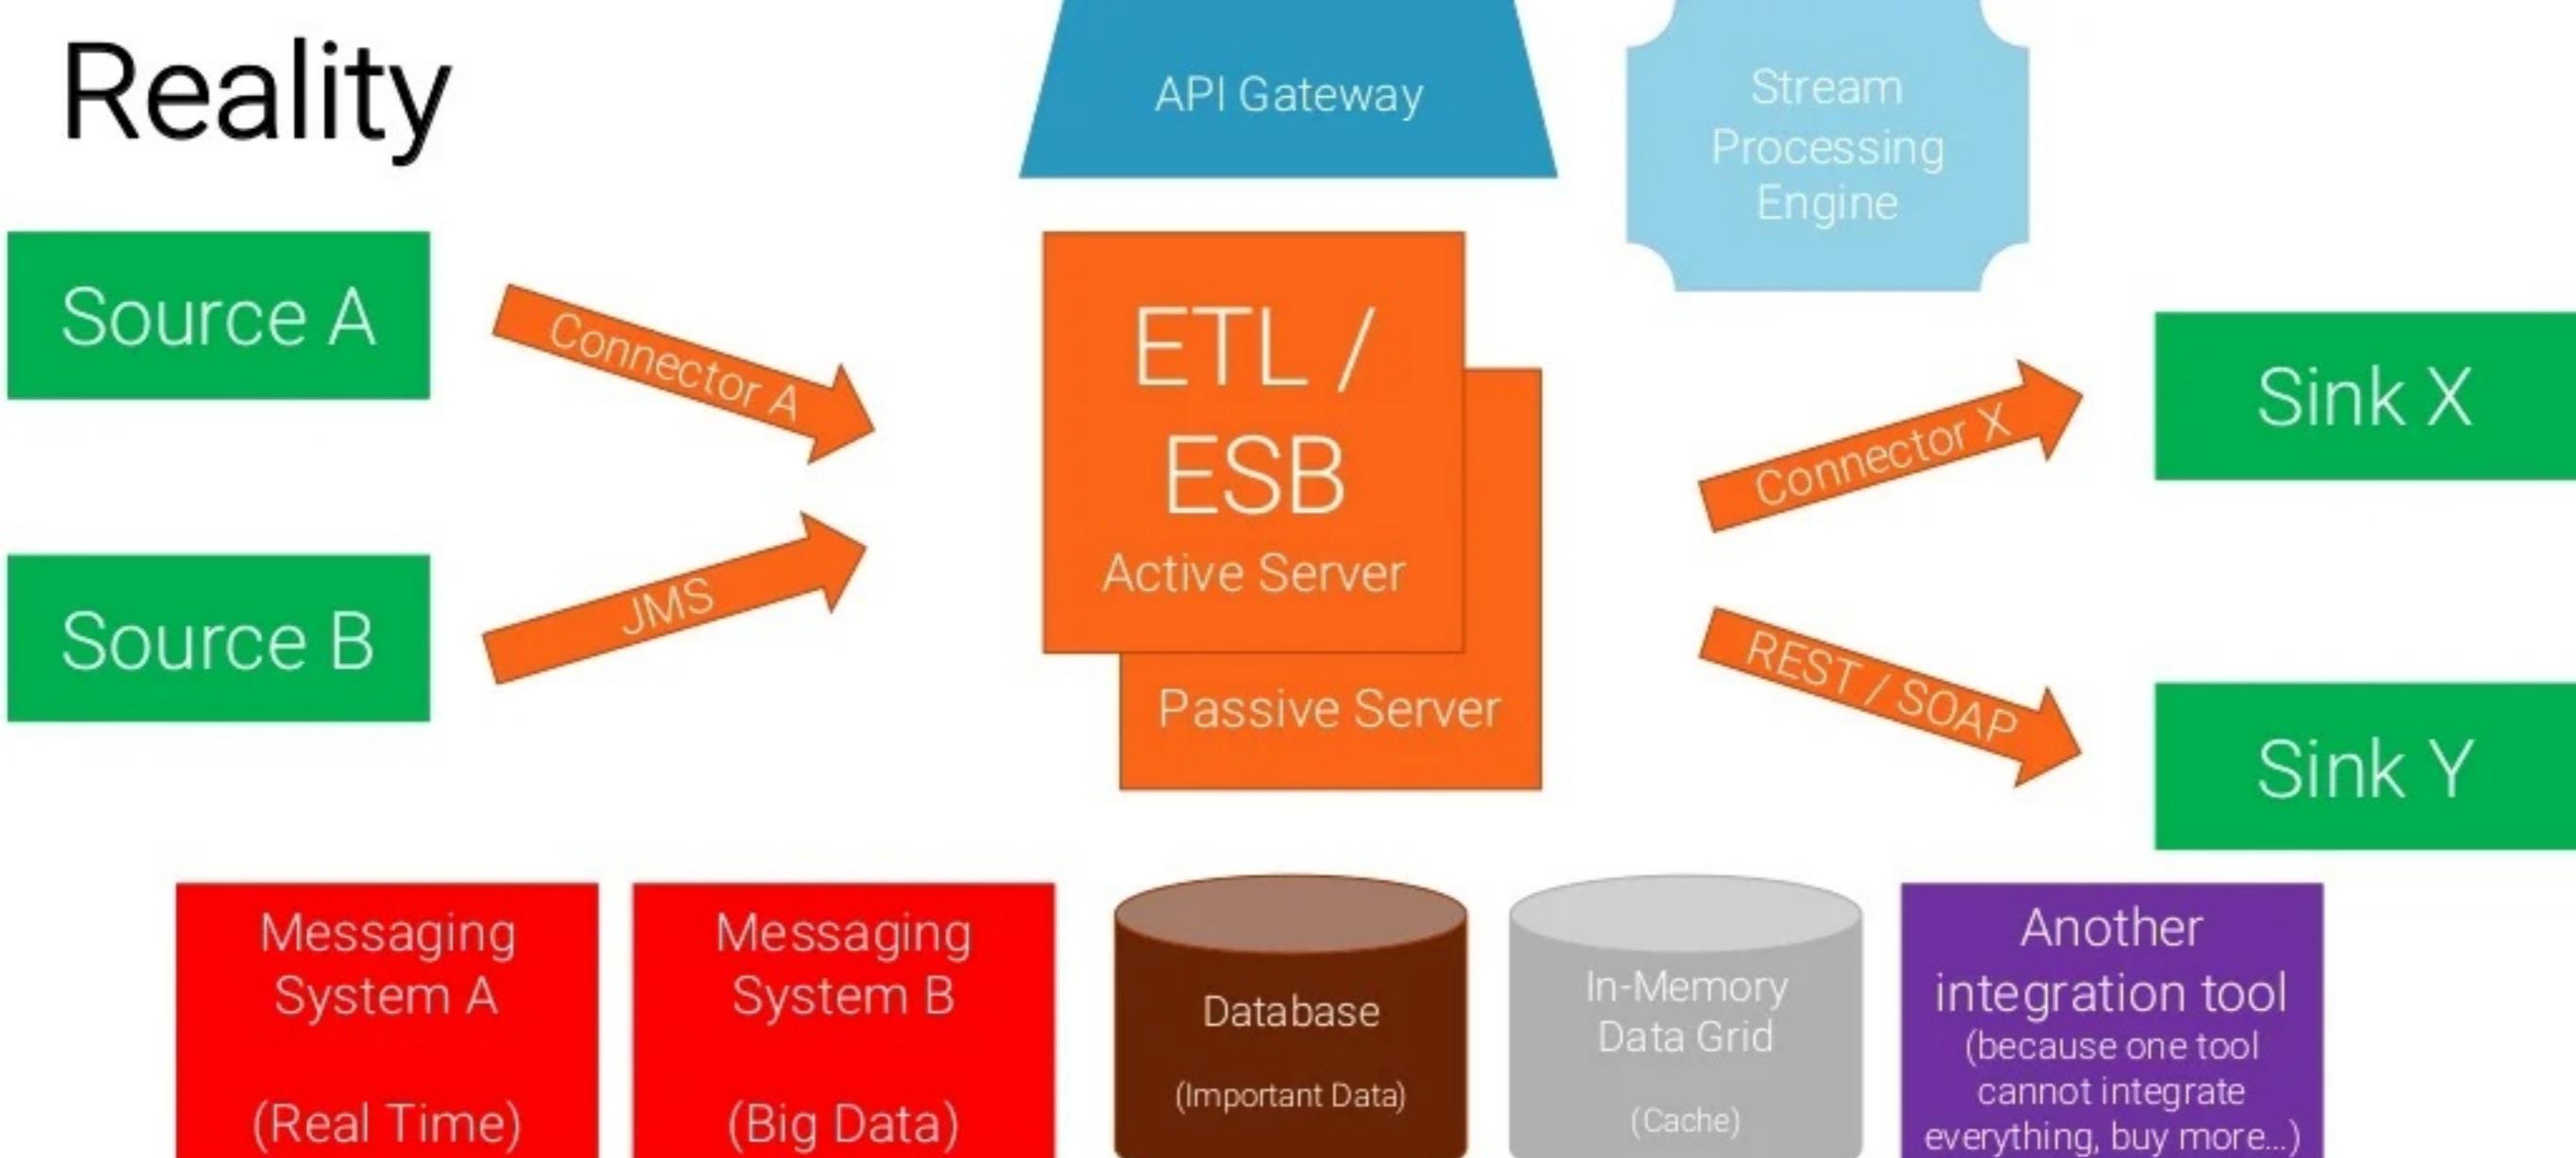
4. Friends

5. Frenemies

# Dream



# Reality



Build your custom integration layer... At least 4-5 different products or open source frameworks

# Challenges with current integration architectures (MQ / ETL / ESB)

---

## Zoo of technologies

Integration platform (Extract-Transform-Load / Enterprise service Bus) + additional "optional" components

Messaging System (Message Queue, Peer-to-Peer)

In-Memory Cache or Data Grid

Database

Streaming Engine

API Gateway

## Architectures with limited scalability and availability

No end-to-end scalability (only parts are built for high volume of messages and high throughput)

No native, built-in scalability (you cannot just add a broker to the running system)

Broker-per-scenario (e.g. hundreds of MQ clusters)

Active-passive clustering

Downtime for maintenance, upgrades, configuration changes

No backwards-compatibility

## Tight coupling

Platform-centric integration implementation (→ no separation of concerns, vendor lock-in)

Synchronous communication via request-response (SOAP / REST / gRPC / ...)

No handling of backpressure or unavailable consumers (push-based messaging queues)

# The World has Changed

---



Mobile



Cloud



Microservices



Internet of Things



Machine Learning

Business Digitalization Trends are Driving the Need to Process  
Events at a whole **new Scale, Speed and Efficiency**



# #NoMQ #NoESB #NoETL can handle this well!

Tight coupling, limited scale, zoo of components

You need to think **event-based** to realize these use cases!

You need to leverage a single, scalable and loosely coupled platform

## Event Streaming Platform!

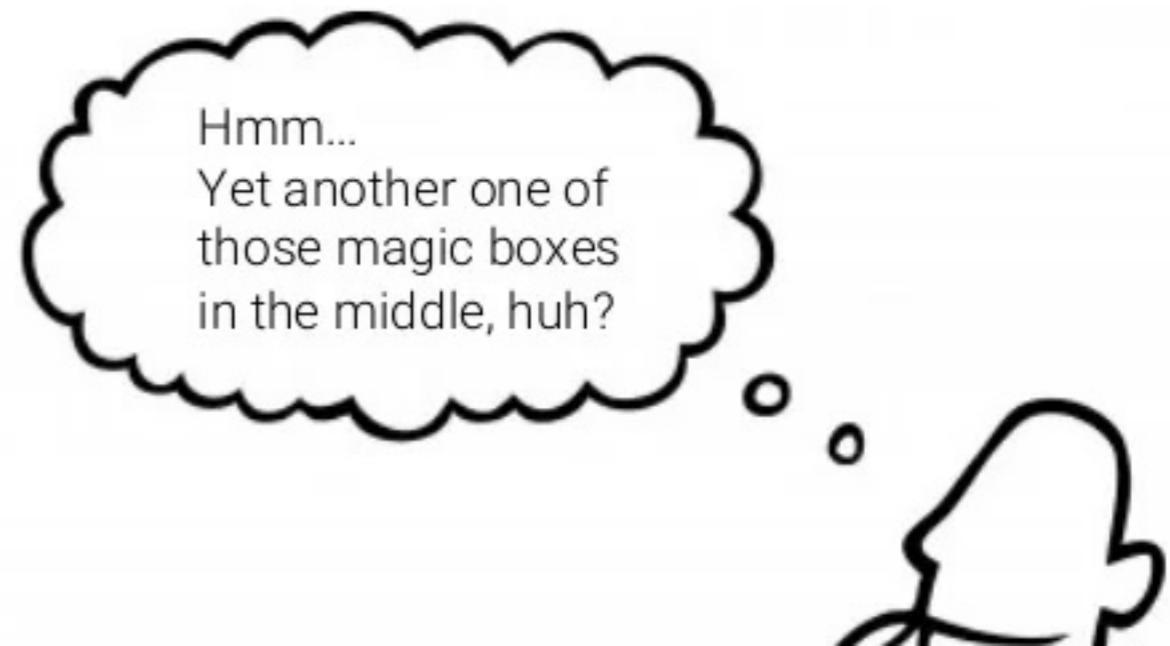
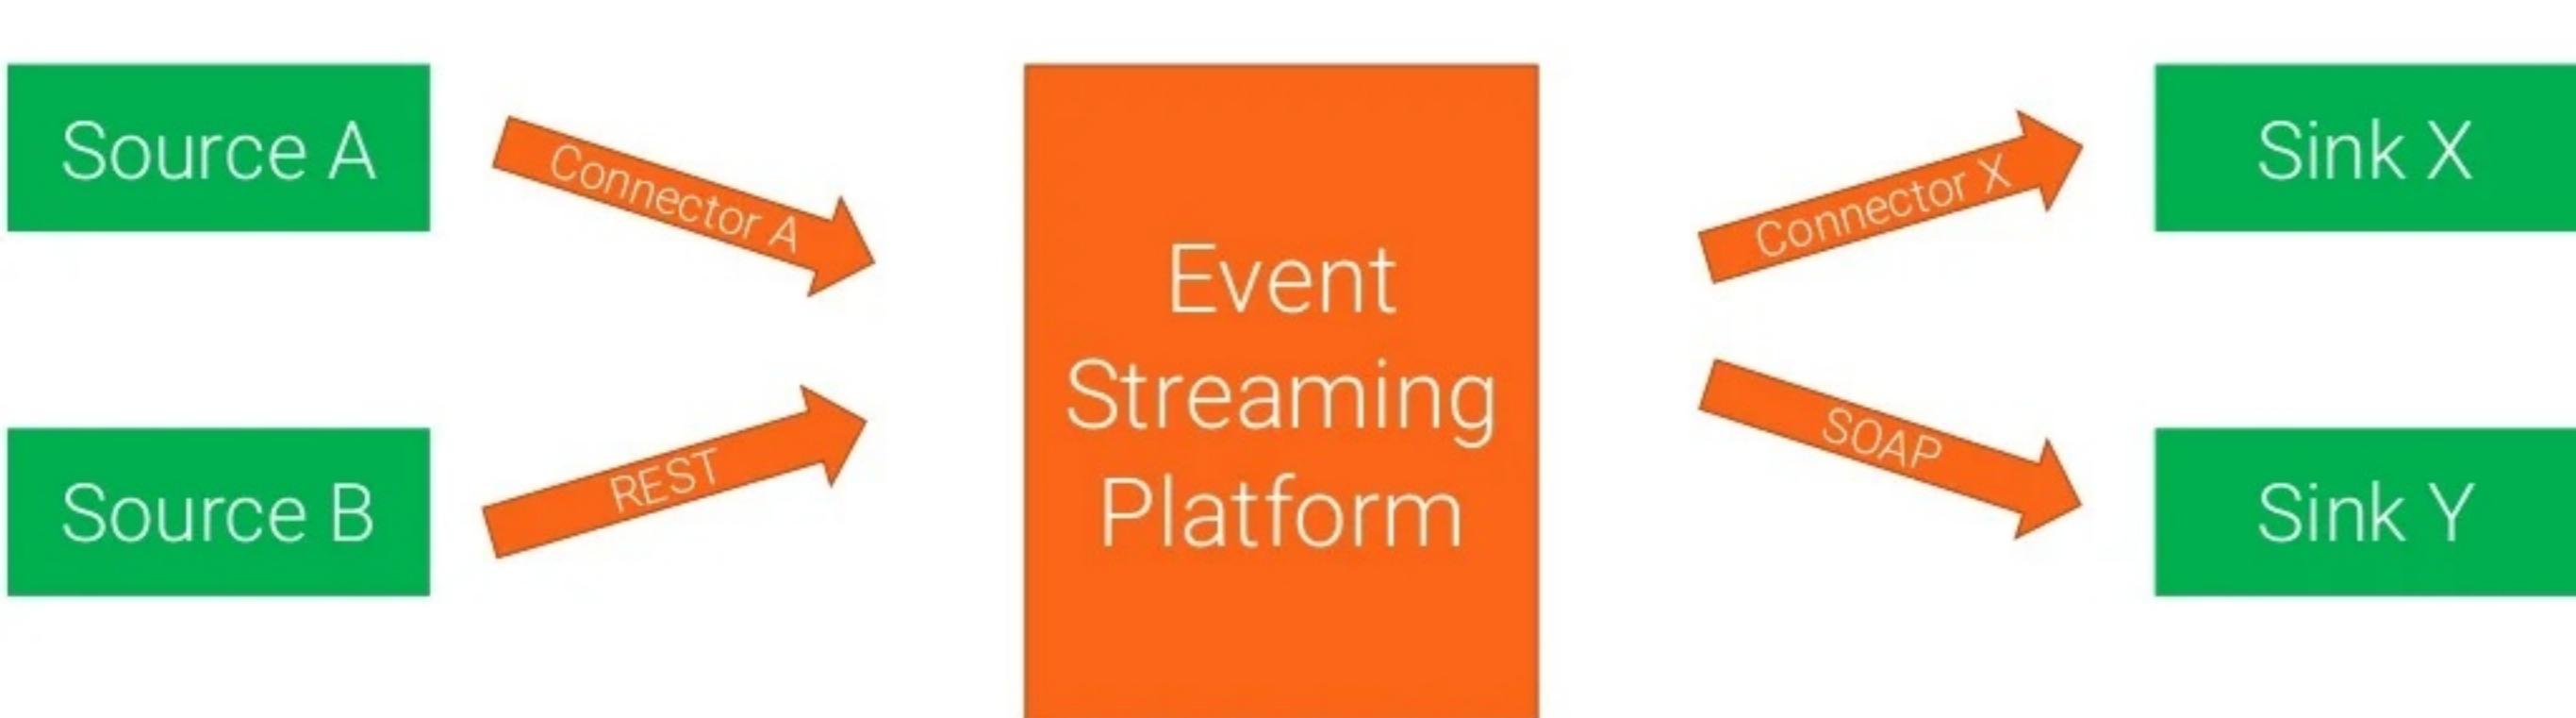
1. Traditional Middleware

**2. Event Streaming Platform**

3. Enemies

4. Friends

5. Frenemies



# Events

What is an event?



Events

# SOMETHING HAPPENED



# Events

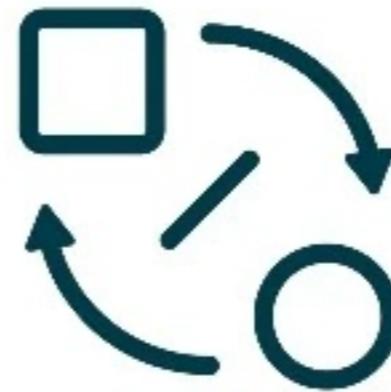
---



A Sale



An Invoice



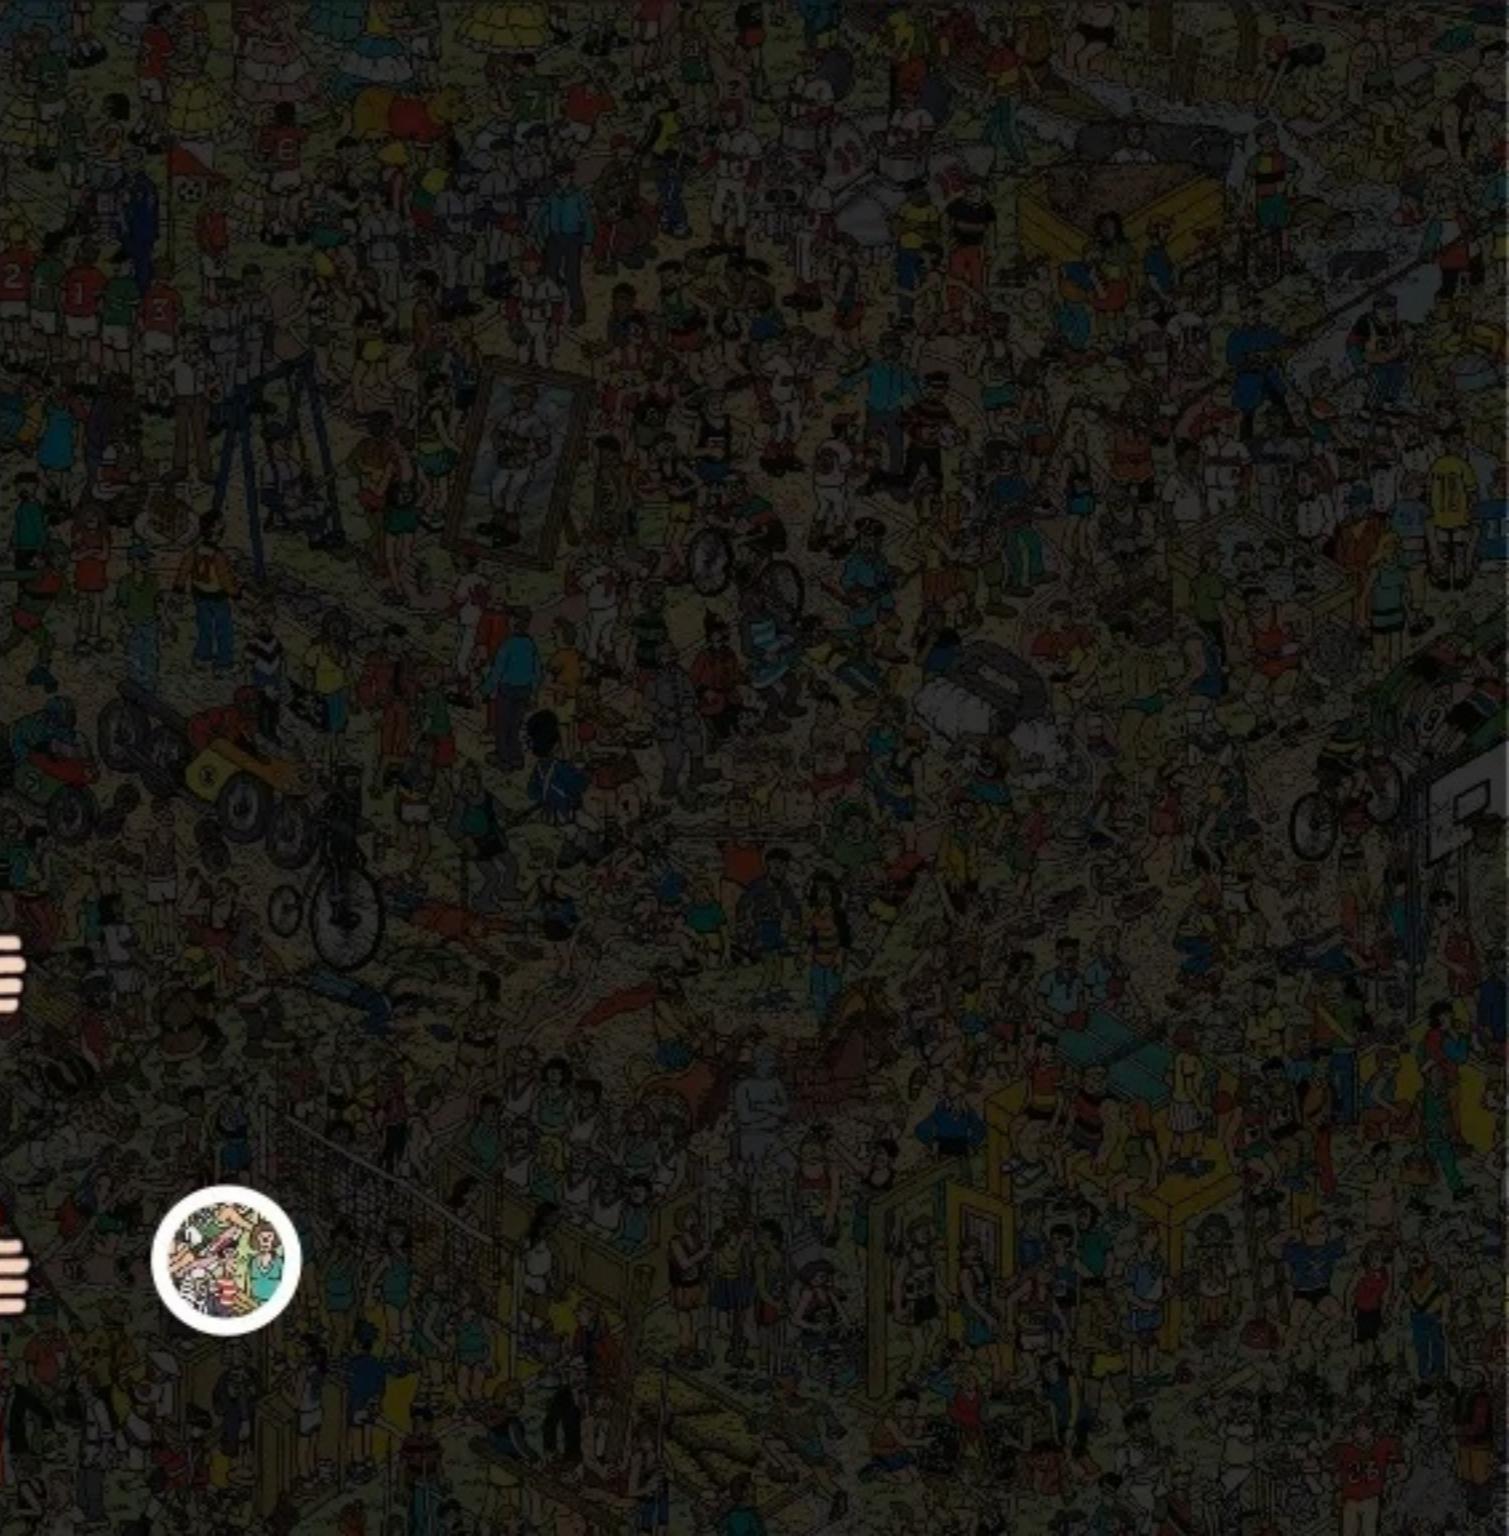
A Trade



A Customer  
Experience

# Where are they?

Events haven't had a  
proper home in  
infrastructure or in code.  
They are implicit.

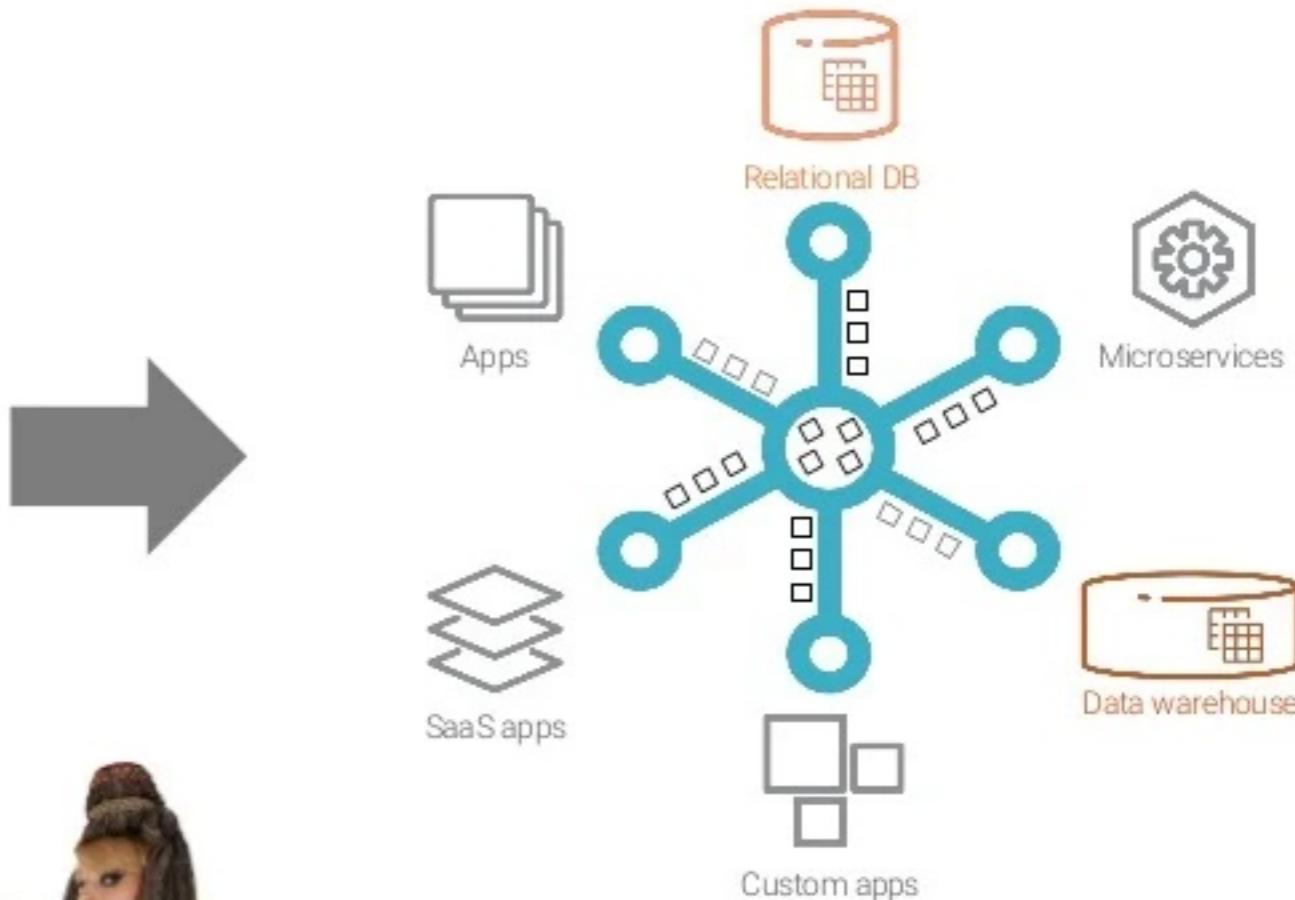


# Implementing an Event-driven Architecture Requires a Paradigm Shift

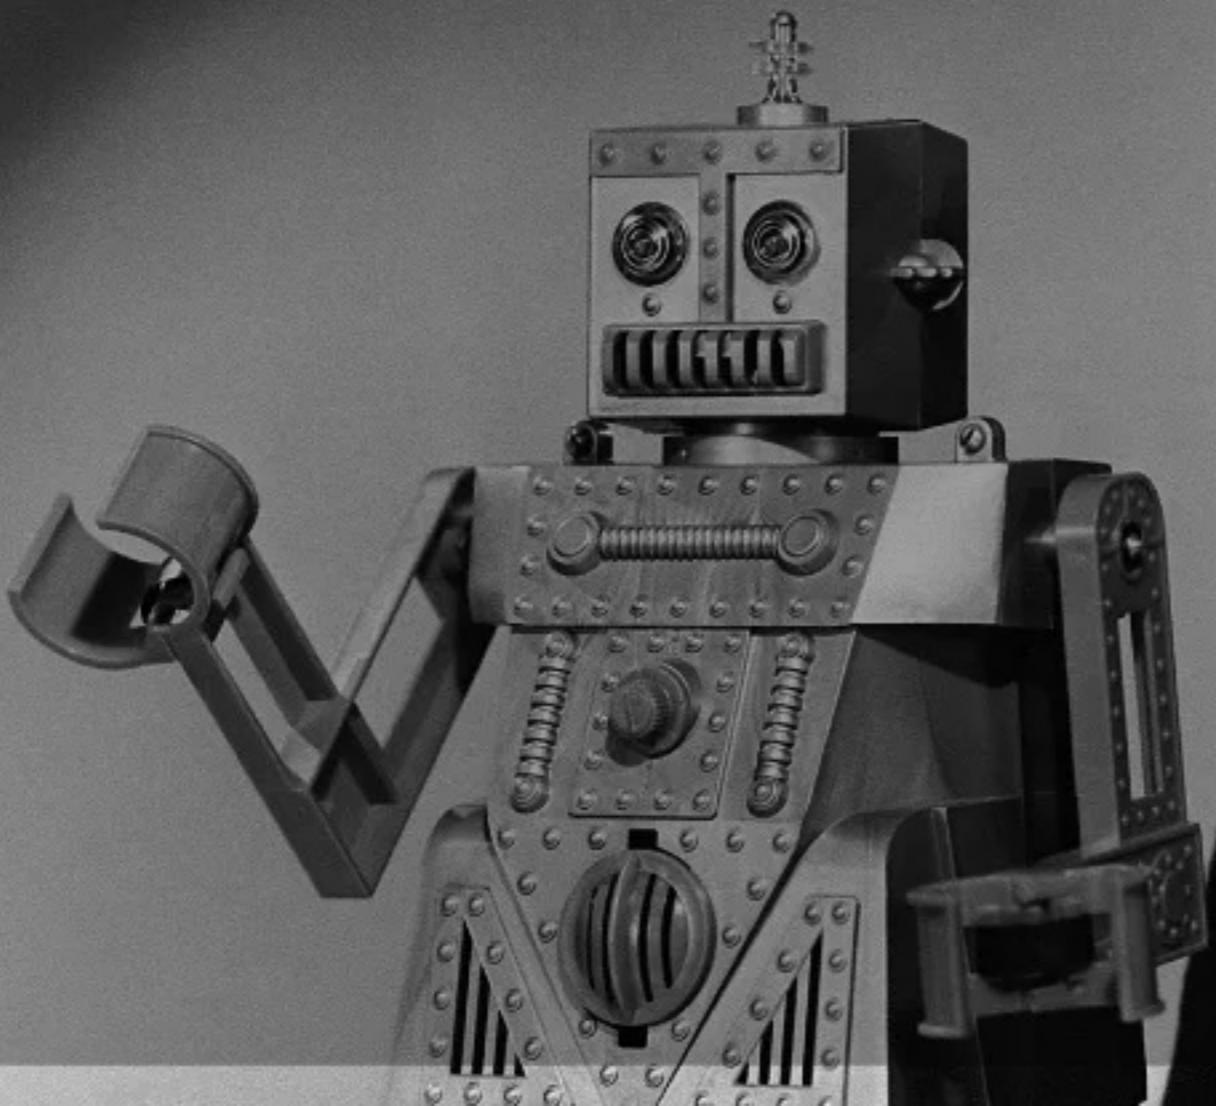
**From data represented in static tables  
and accessed with RPC...**

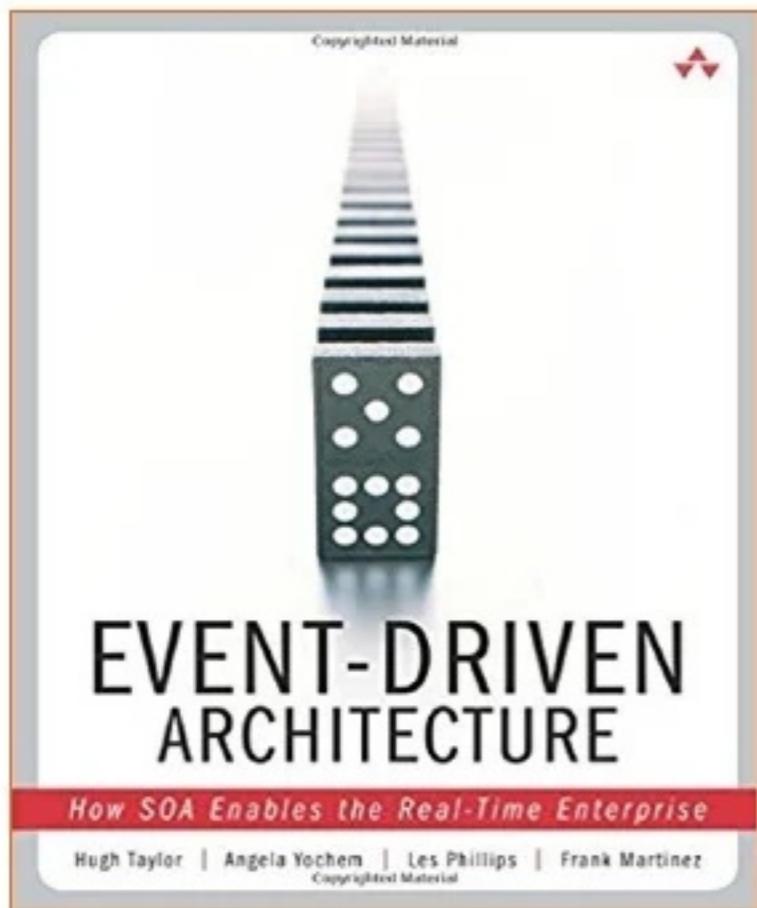


**...to data represented as streams of events**

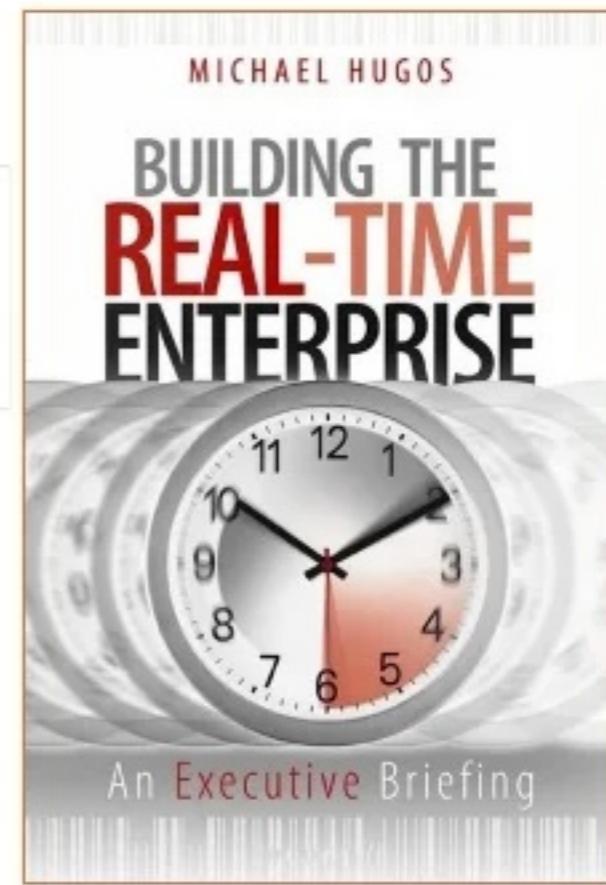


Haven't we seen all  
this before?





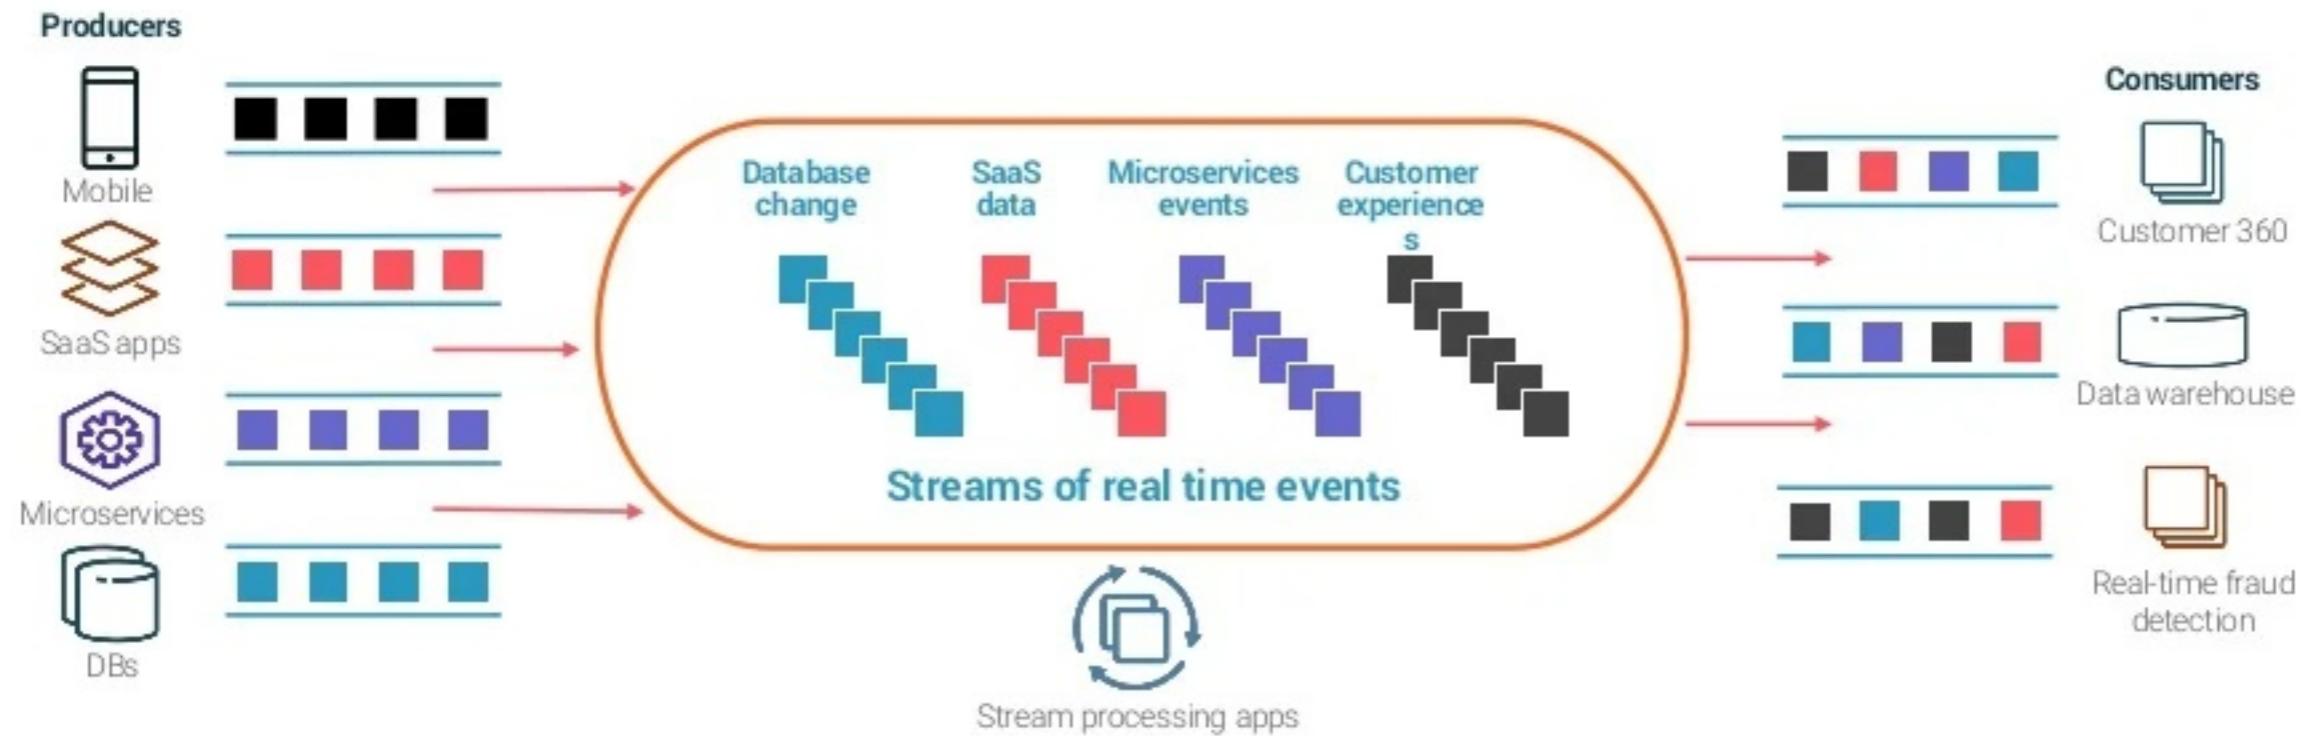
(Published in 2009)



(Published in 2004)

# What's different this time around?

# A Streaming Platform is the Underpinning of an Event-driven Architecture



## Ubiquitous connectivity

Globally scalable platform for all event producers and consumers

## Immediate data access

Data accessible to all consumers in real time

## Single System of record

Persistent storage to enable reprocessing of past events

## Continuous queries

Stream processing capabilities for in-line data transformation<sup>19</sup>



# Agenda

- Traditional Middleware
- Event Streaming Platform
- **Enemies**
- Friends
- Frenemies



Middleware: Message Queue... ETL... Enterprise Service Bus...  
Event Streaming Platform: Apache Kafka.  
Spoilt for Choice!

Why do you need an Event Streaming Platform? Remember:

---

# The World has Changed!



Mobile



Cloud



Microservices



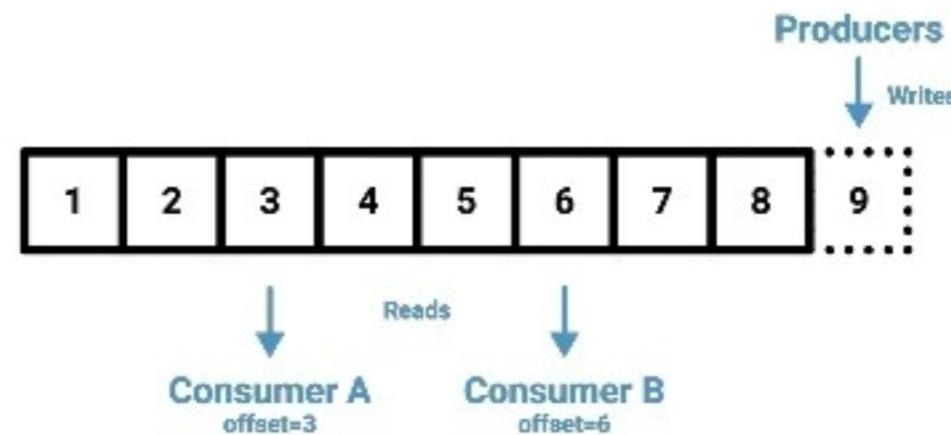
Internet of Things



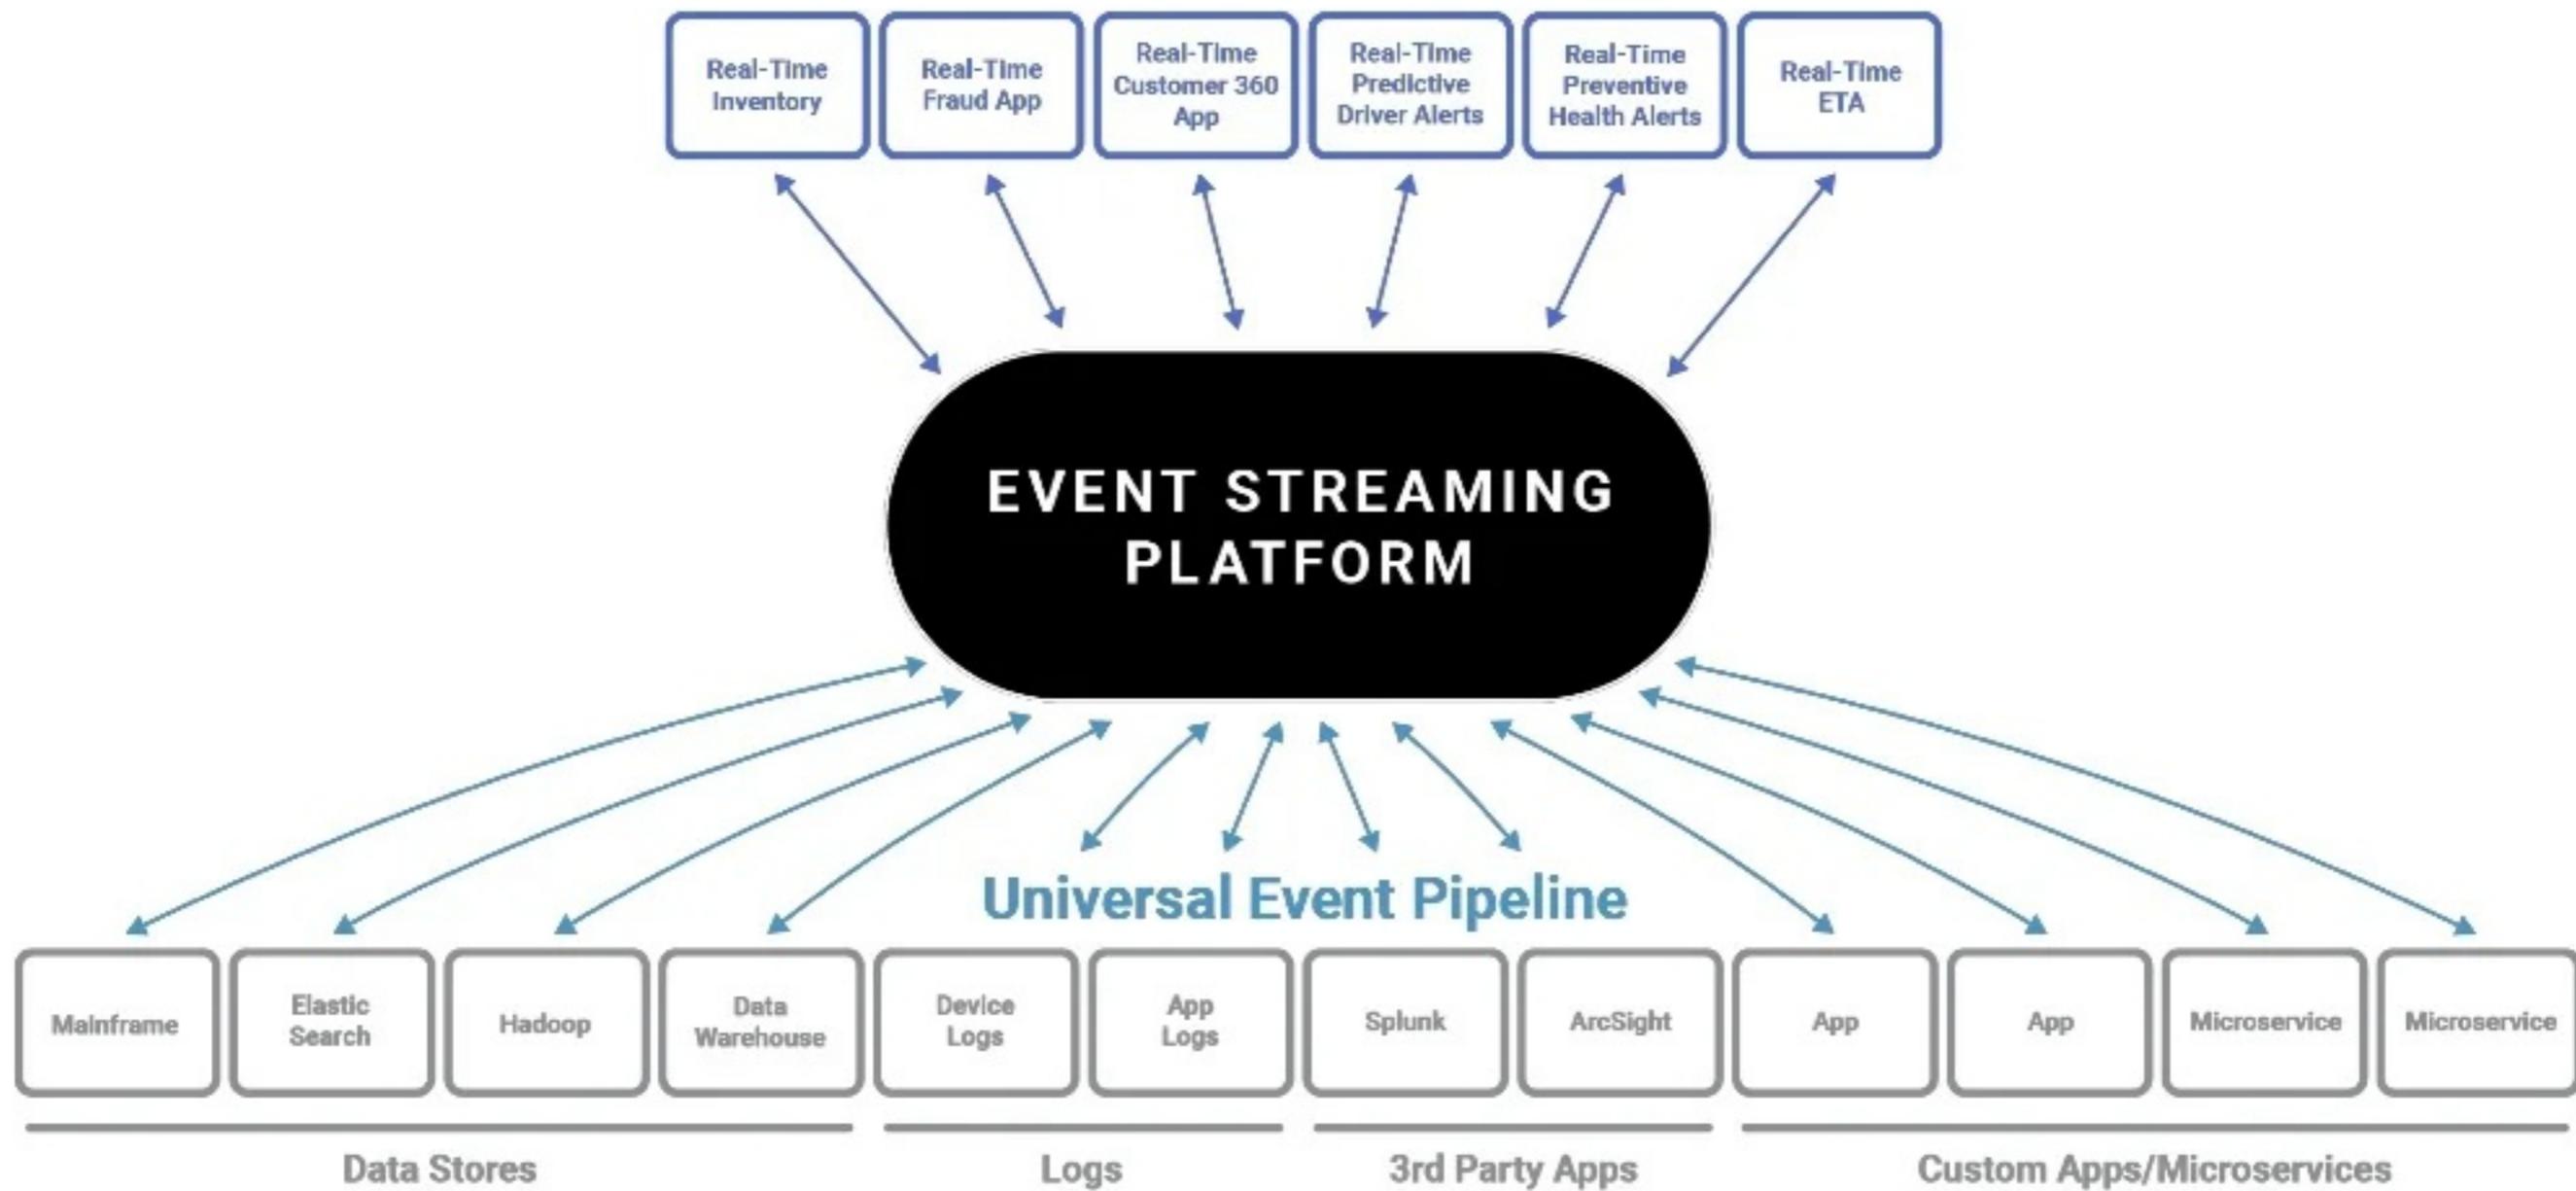
Machine Learning

Business Digitalization Trends are Driving the Need to Process Events at a whole **new Scale, Speed and Efficiency**

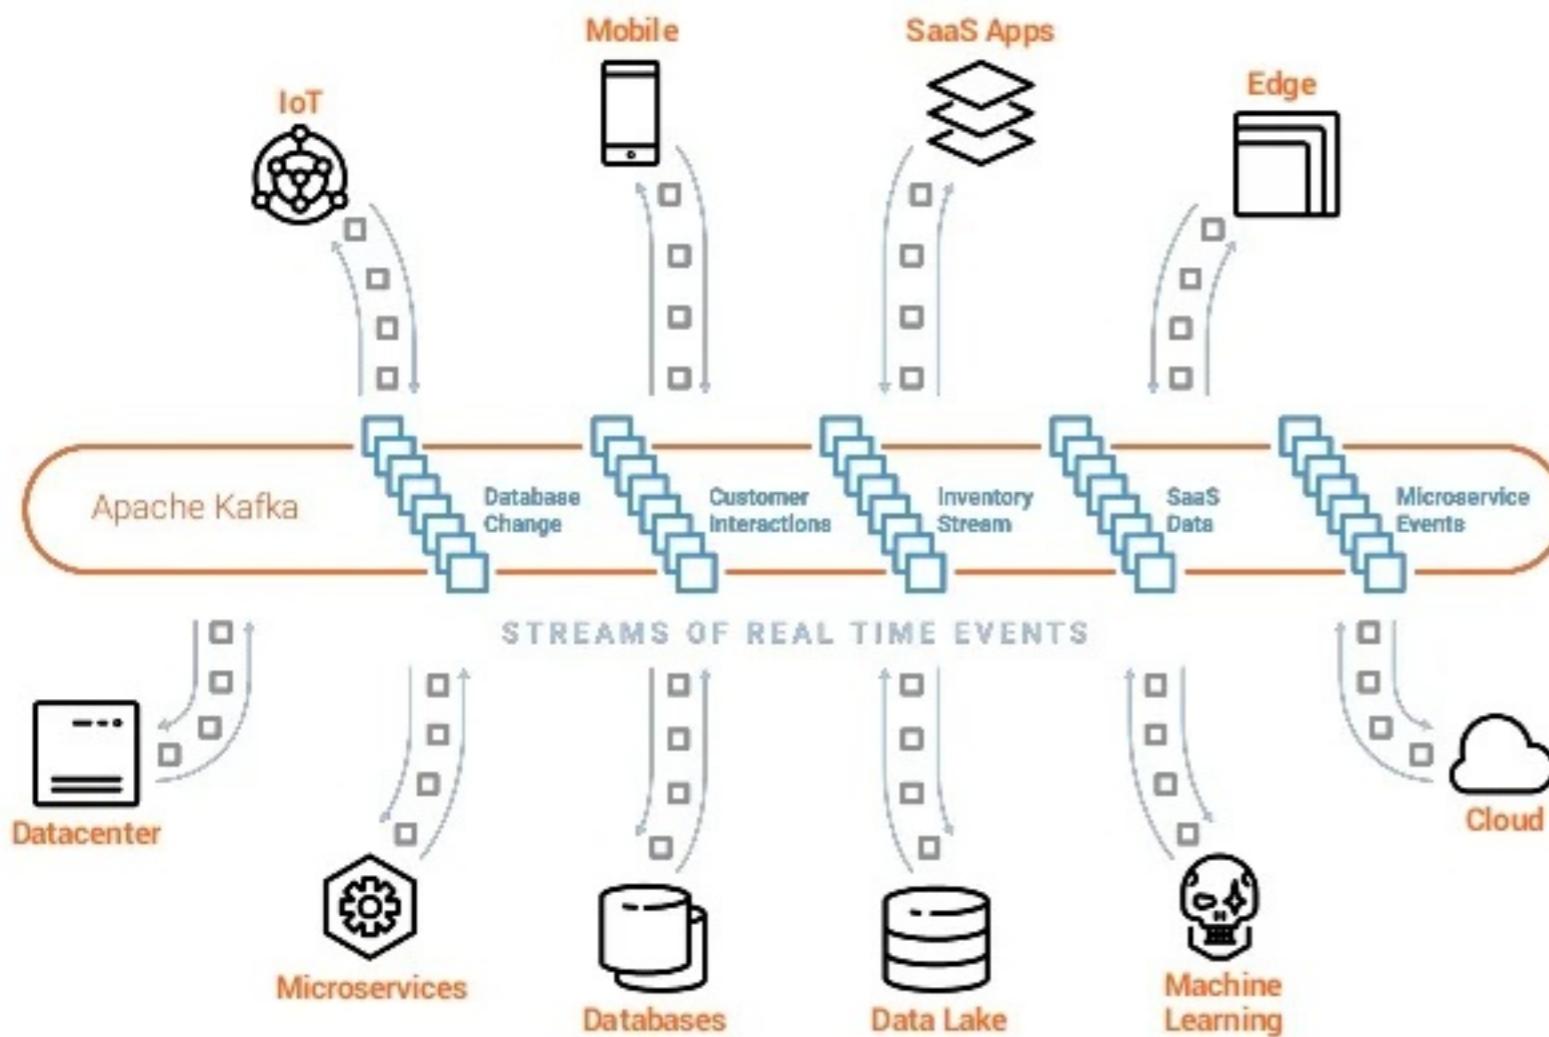
# Event Streaming Paradigm



# Contextual Event-Driven Apps



# Apache Kafka: The De-facto Standard for Real-Time Event Streaming



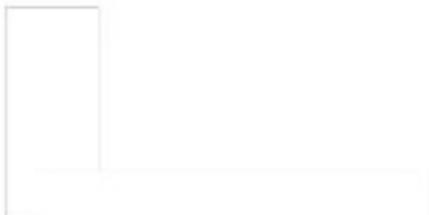
- Global-scale
- Real-time
- Persistent Storage
- Stream Processing



> 4.5 trillion messages / day

The Netflix logo, consisting of the word "NETFLIX" in large red block letters.

> 6 Petabytes / day



The Pinterest logo, consisting of the word "Pinterest" in red cursive script. To its right is the slogan "You name it" in a smaller black sans-serif font.

# Event Streaming Platform - Value per Use Case



Business Value



# Why Apache Kafka instead of traditional middleware?

## Event Streaming Platform

- The core is event-based
- Supports real time stream processing
- But also Fire-and-Forget, Publish / Subscribe, Request-Response / RPC, Batch, ...

## Single infrastructure

- Messaging, storage, processing
- Extreme Scale and throughput

## Reliability and zero downtime (“built for failure”)

- High availability
- Rolling upgrades and dynamic configuration changes
- Backwards compatibility

## Decoupling of clients

- Agile microservices
- Dumb pipes, smart endpoints
- Handling backpressure

## No vendor lock-in



## Why Apache Kafka instead of traditional middleware?

---



**"Eat your own dog good!"**

Enemies! .... More components, clusters, technologies means more conflicts, incompatibility, operations burden!



VS.



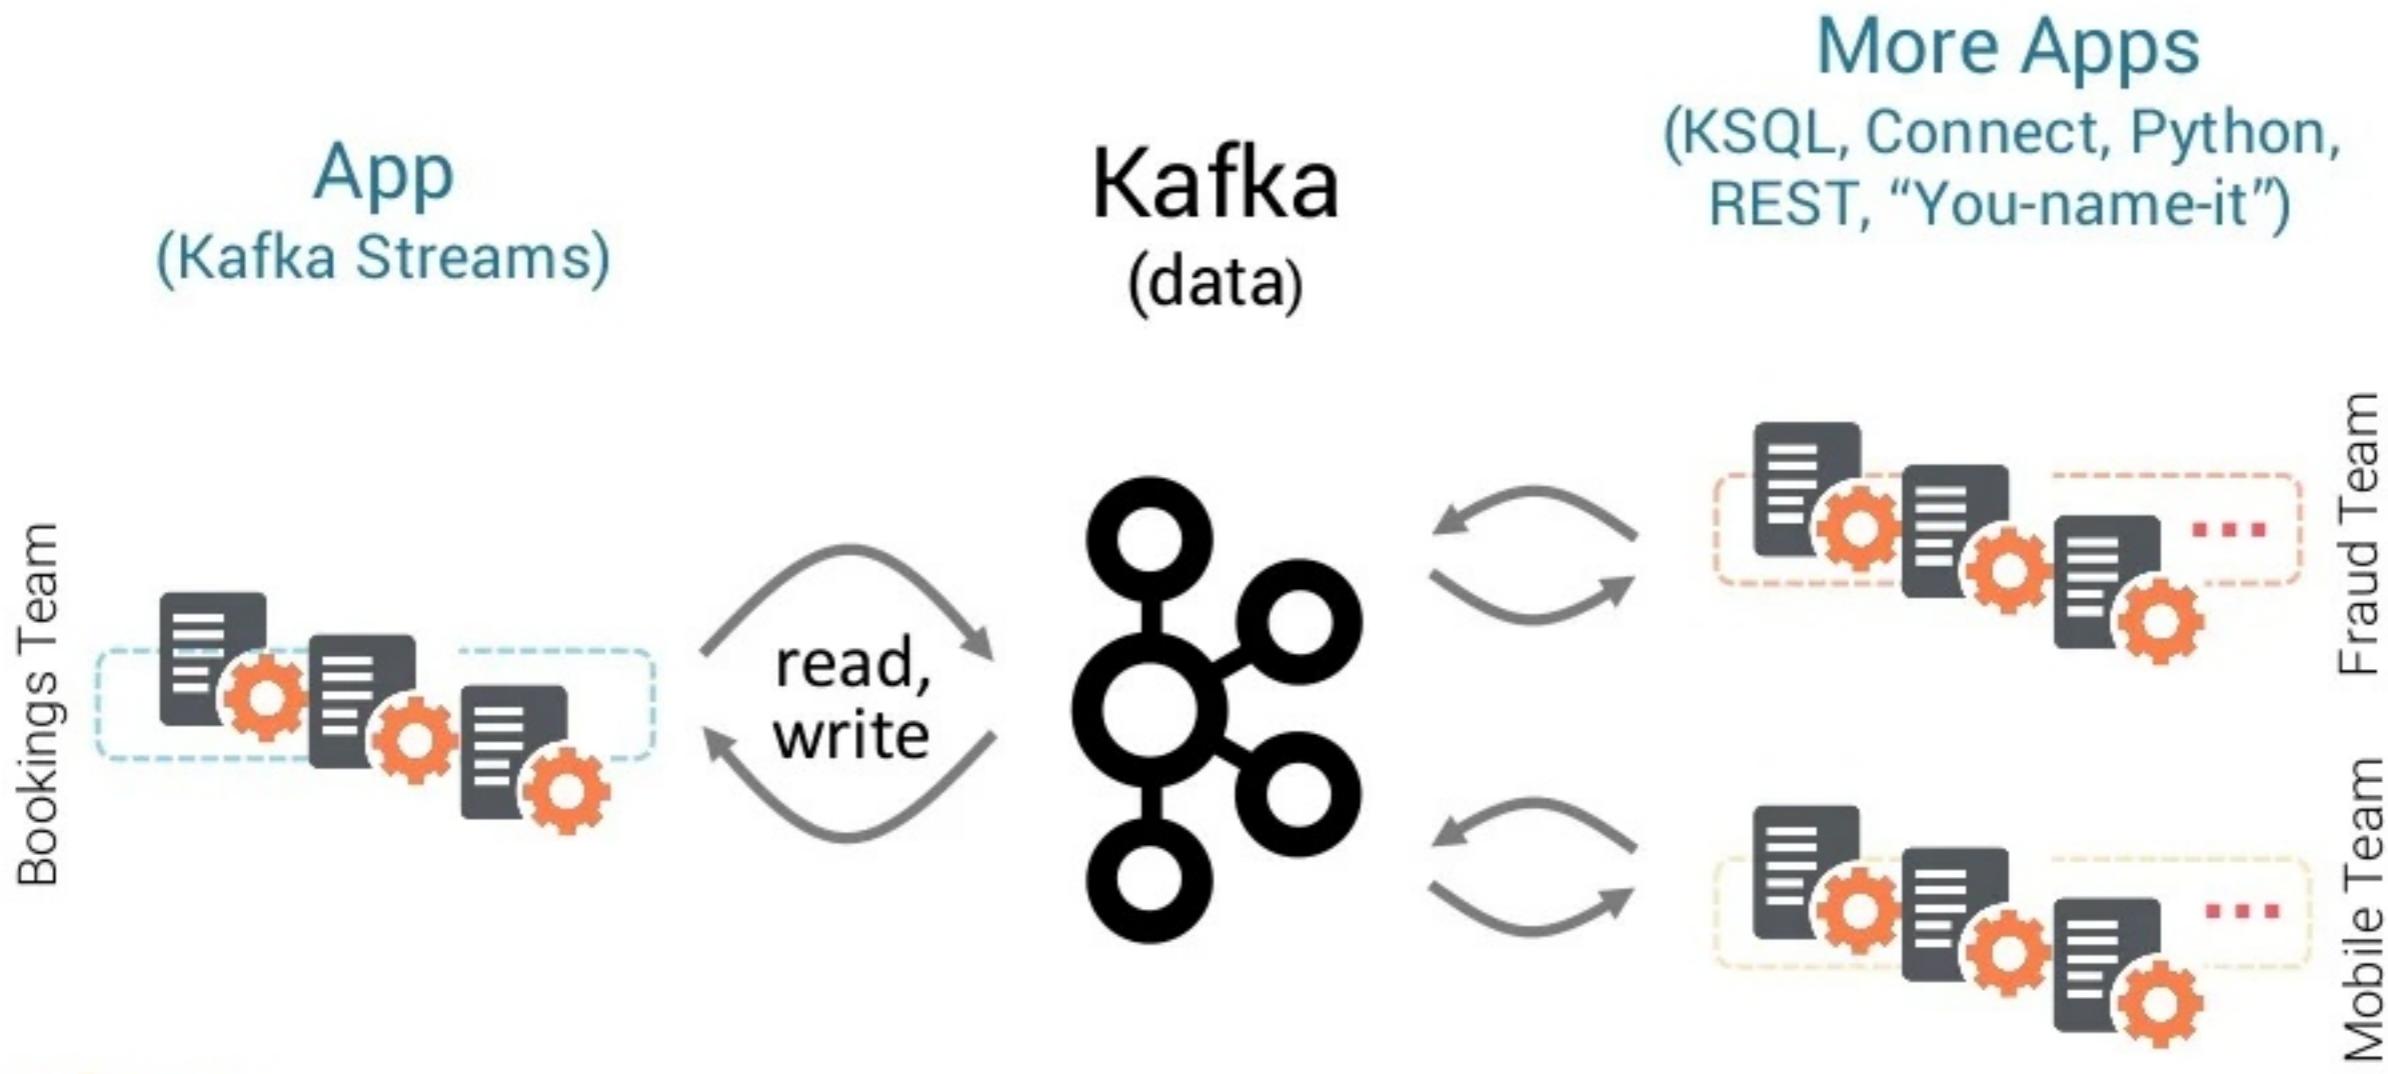
Messaging: Kafka Core  
Storage: Kafka Core  
Caching: Kafka Core  
Real-Time, Batch: Kafka Clients  
Integration: Kafka Connect  
Stream Processing: Kafka Streams / KSQL  
Request-Response: REST Proxy  
**"Eat your own dog good"**



No need for another infrastructure, cluster, database... High availability and scale handled by Kafka Topics!

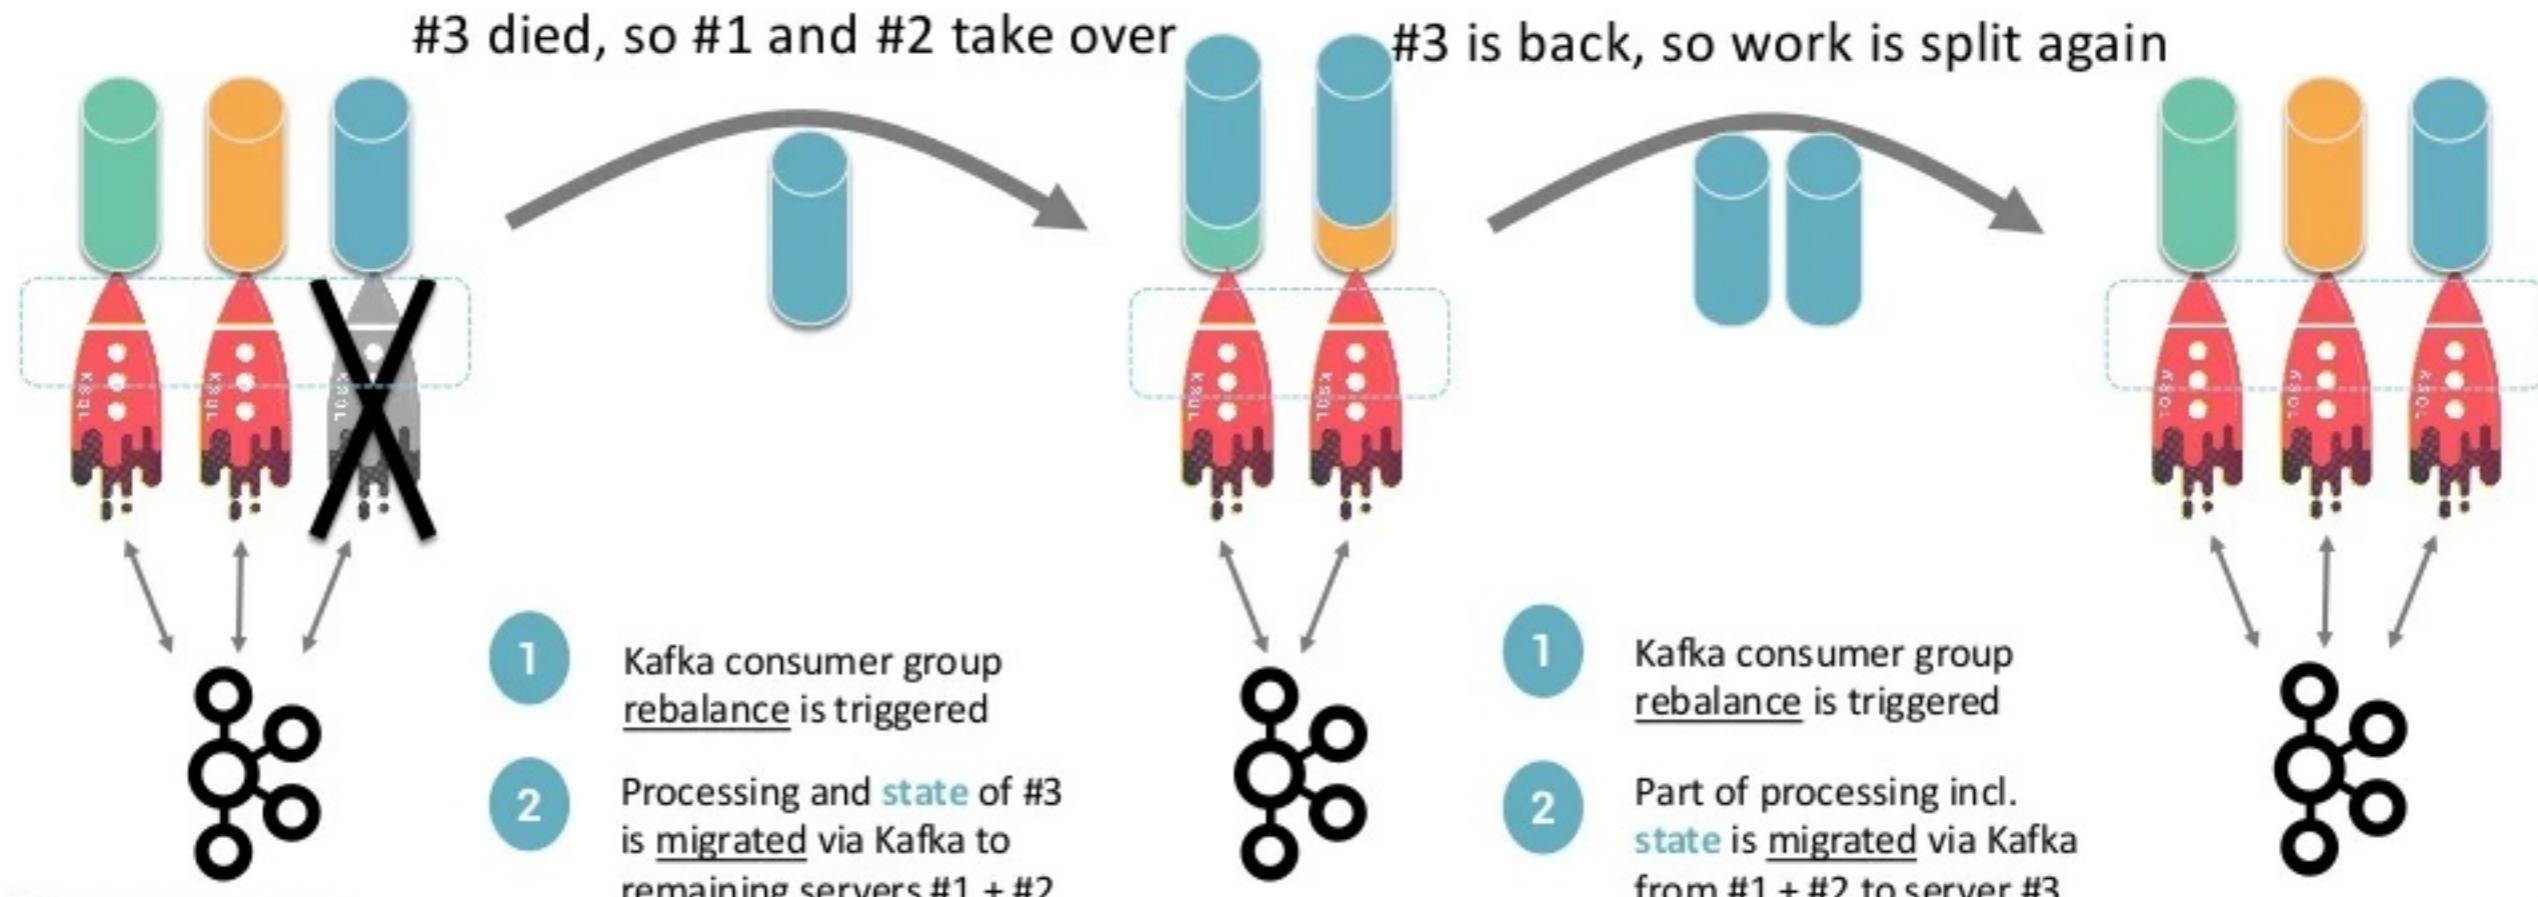
Independent, scalable, reliable components

---

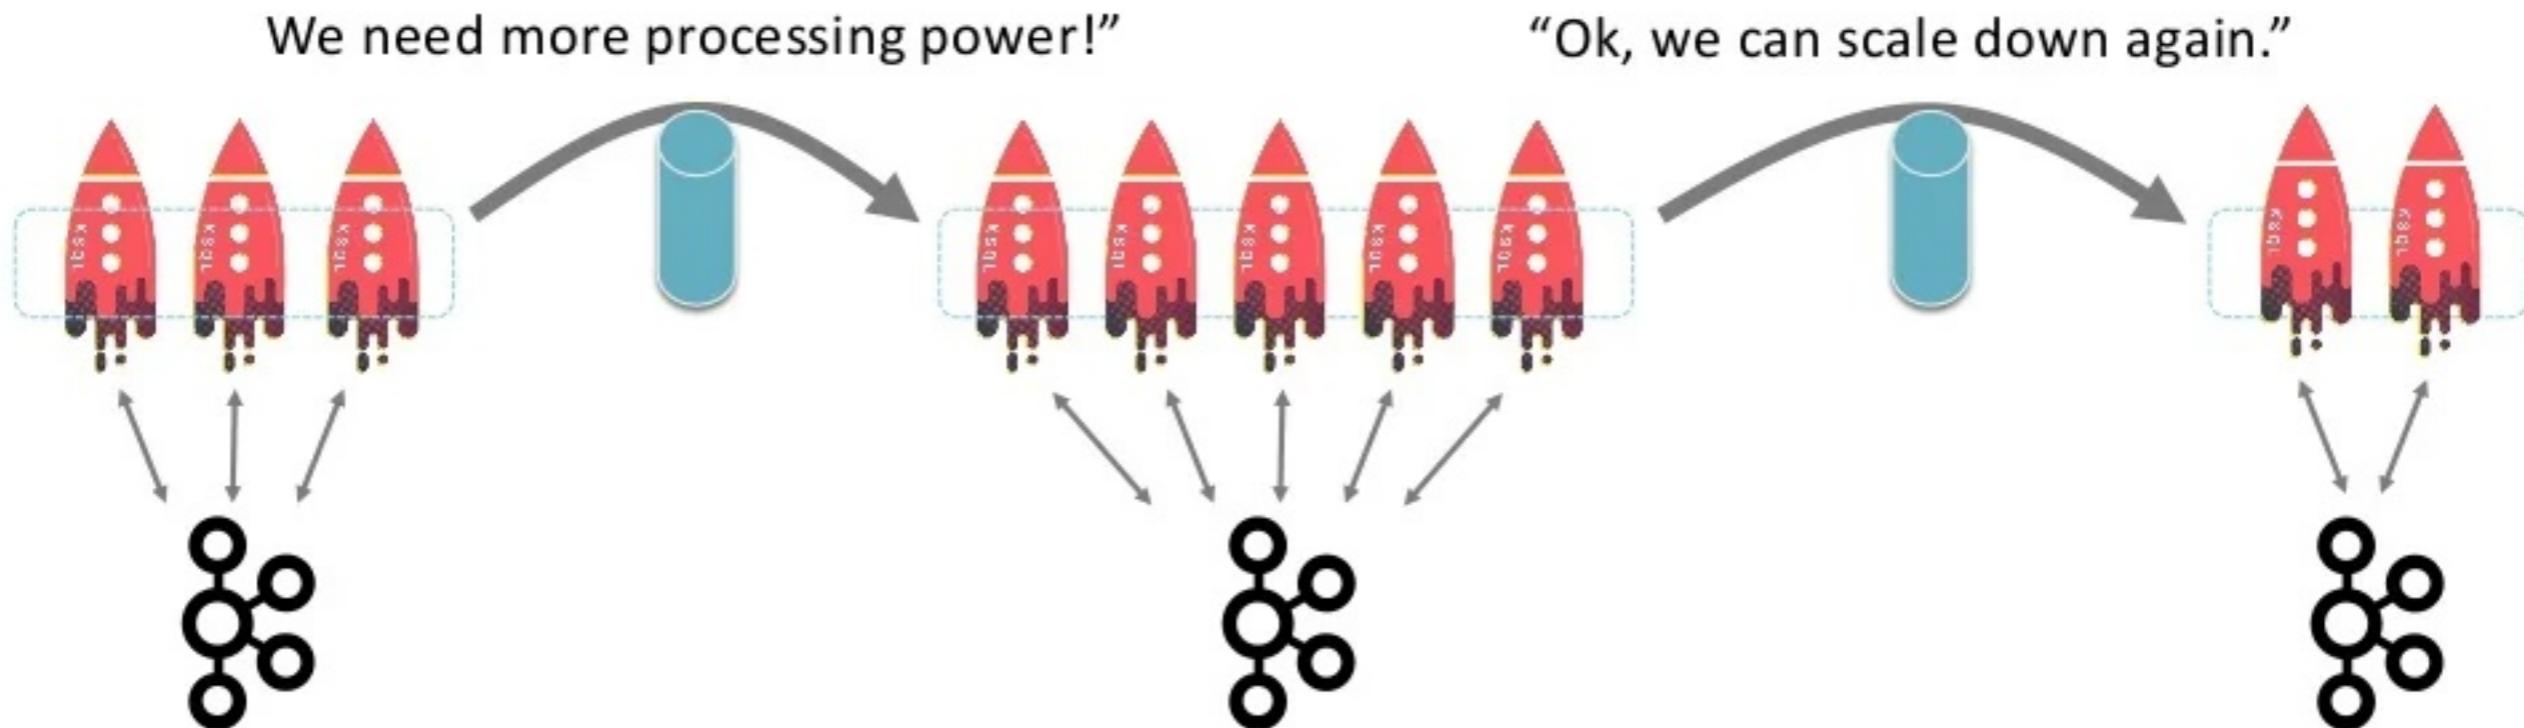


# Example - Kafka Streams / KSQL : Fault-tolerance, powered by Kafka

Processing fails over automatically, without data loss or miscomputation.



You can add, remove, restart servers during live operations.



# Don't use your ESB knowledge with Kafka!

## Recreating ESB antipatterns with Kafka

MAY  
2018

HOLD ?

Kafka is becoming very popular as a messaging solution, and along with it, [Kafka Streams](#) is at the forefront of the wave of interest in streaming architectures. Unfortunately, as they start to embed Kafka at the heart of their data and application platforms, we're seeing some organizations **recreating ESB antipatterns with Kafka** by centralizing the Kafka ecosystem components — such as connectors and stream processors — instead of allowing these components to live with product or service teams. This reminds us of seriously problematic ESB antipatterns, where more and more logic, orchestration and transformation were thrust into a centrally managed ESB, creating a significant dependency on a centralized team. We're calling this out to dissuade further implementations of this flawed pattern.

**ThoughtWorks®**



# Agenda

- Traditional Middleware
- Event Streaming Platform
- Enemies
- **Friends**
- Frenemies



## Plenty of integration options between Kafka and traditional middleware



- Kafka Connect connectors  
(JMS, IBM MQ, RabbitMQ, etc.)
- JMS Client  
(Kafka-native JMS Implementation)
- ESB or ETL tools  
with their own connectors
- Kafka's Client APIs  
(like Java, .NET, Go, Python, Javascript)
- REST Proxy
- Etc.

# Why friends?

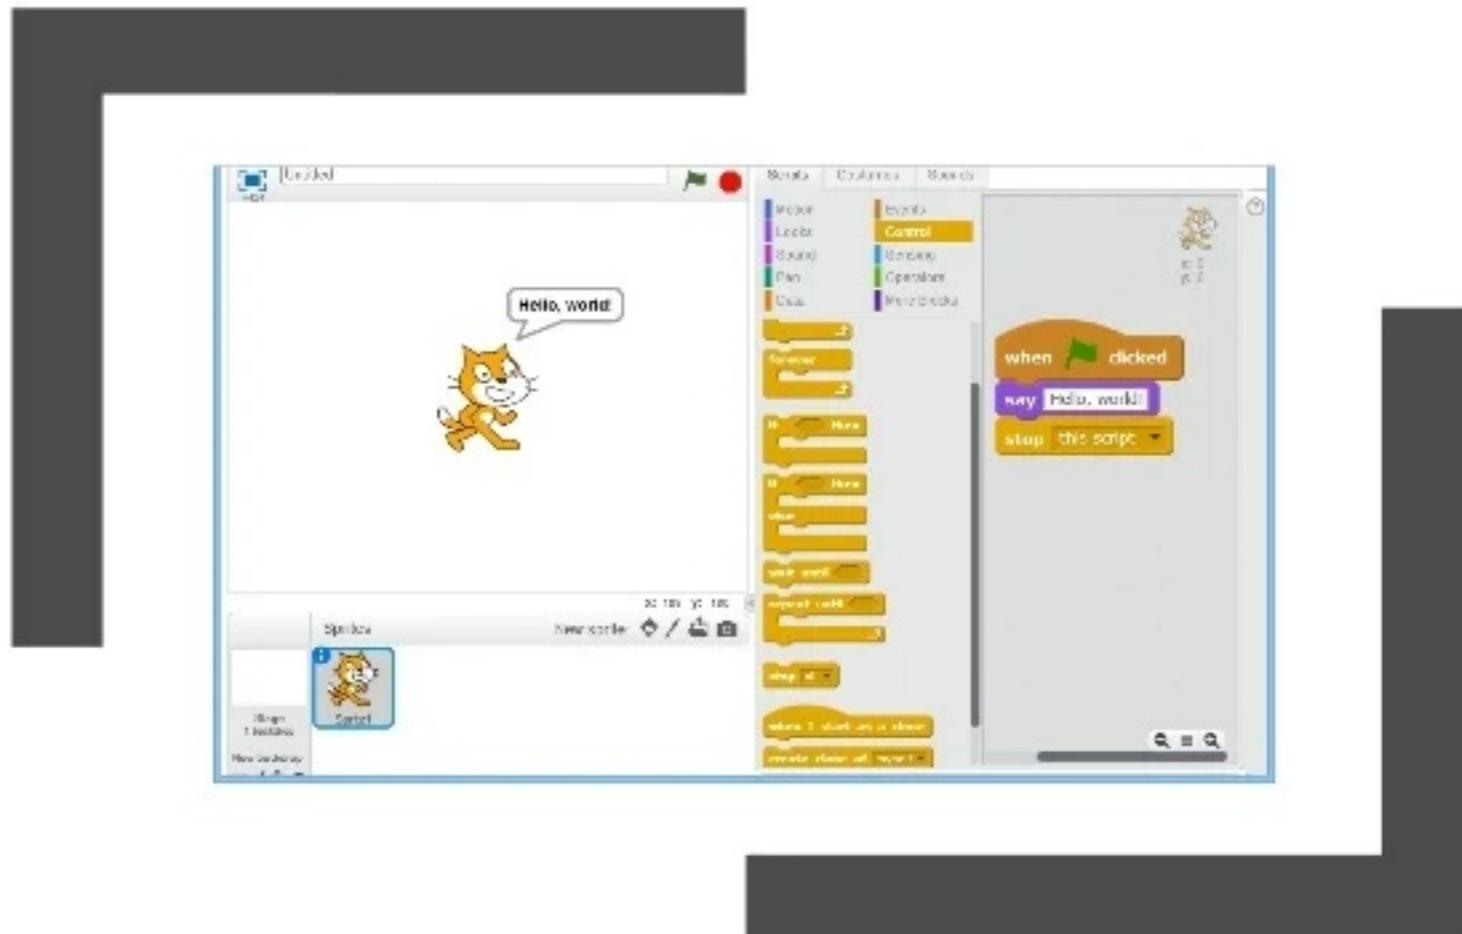
Kafka is NOT THE ALLROUNDER for every single problem!

Maybe you don't need a scalable, reliable, distributed system, but "just":

- Integration with legacy components (Cobol, Edifact, ...)
- Point-to-point messaging with active / passive HA for ("a few") messages
- Very specific messaging solution for "less than millisecond performance"
- API Management
- "Real" Batch Processing (Hadoop, Flink, Informatica, ...)
- ...



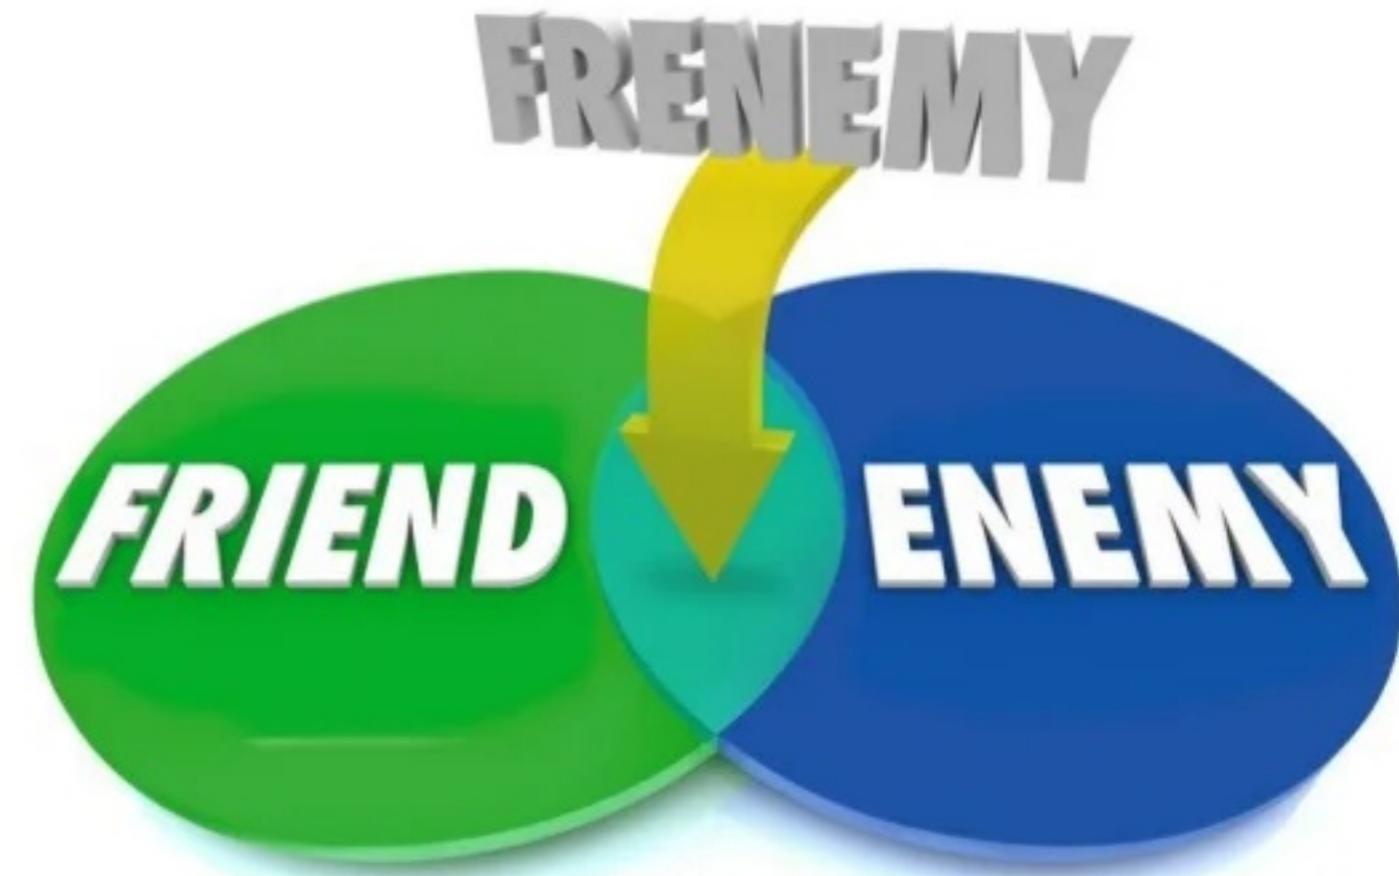
# Choose the right integration technology



- Visual coding for complex graphical mappings
- Integration components for complex legacy standard software and protocols
- (SAP BAPI / iDoc, Cobol, Edifact, SOAP / WS\*, etc.)

## Agenda

- Traditional Middleware
- Event Streaming Platform
- Enemies
- Friends
- **Frenemies**



## Big Bang fails for critical 24/7 deployments

---



No!

Some technologies never die... Cough...

---

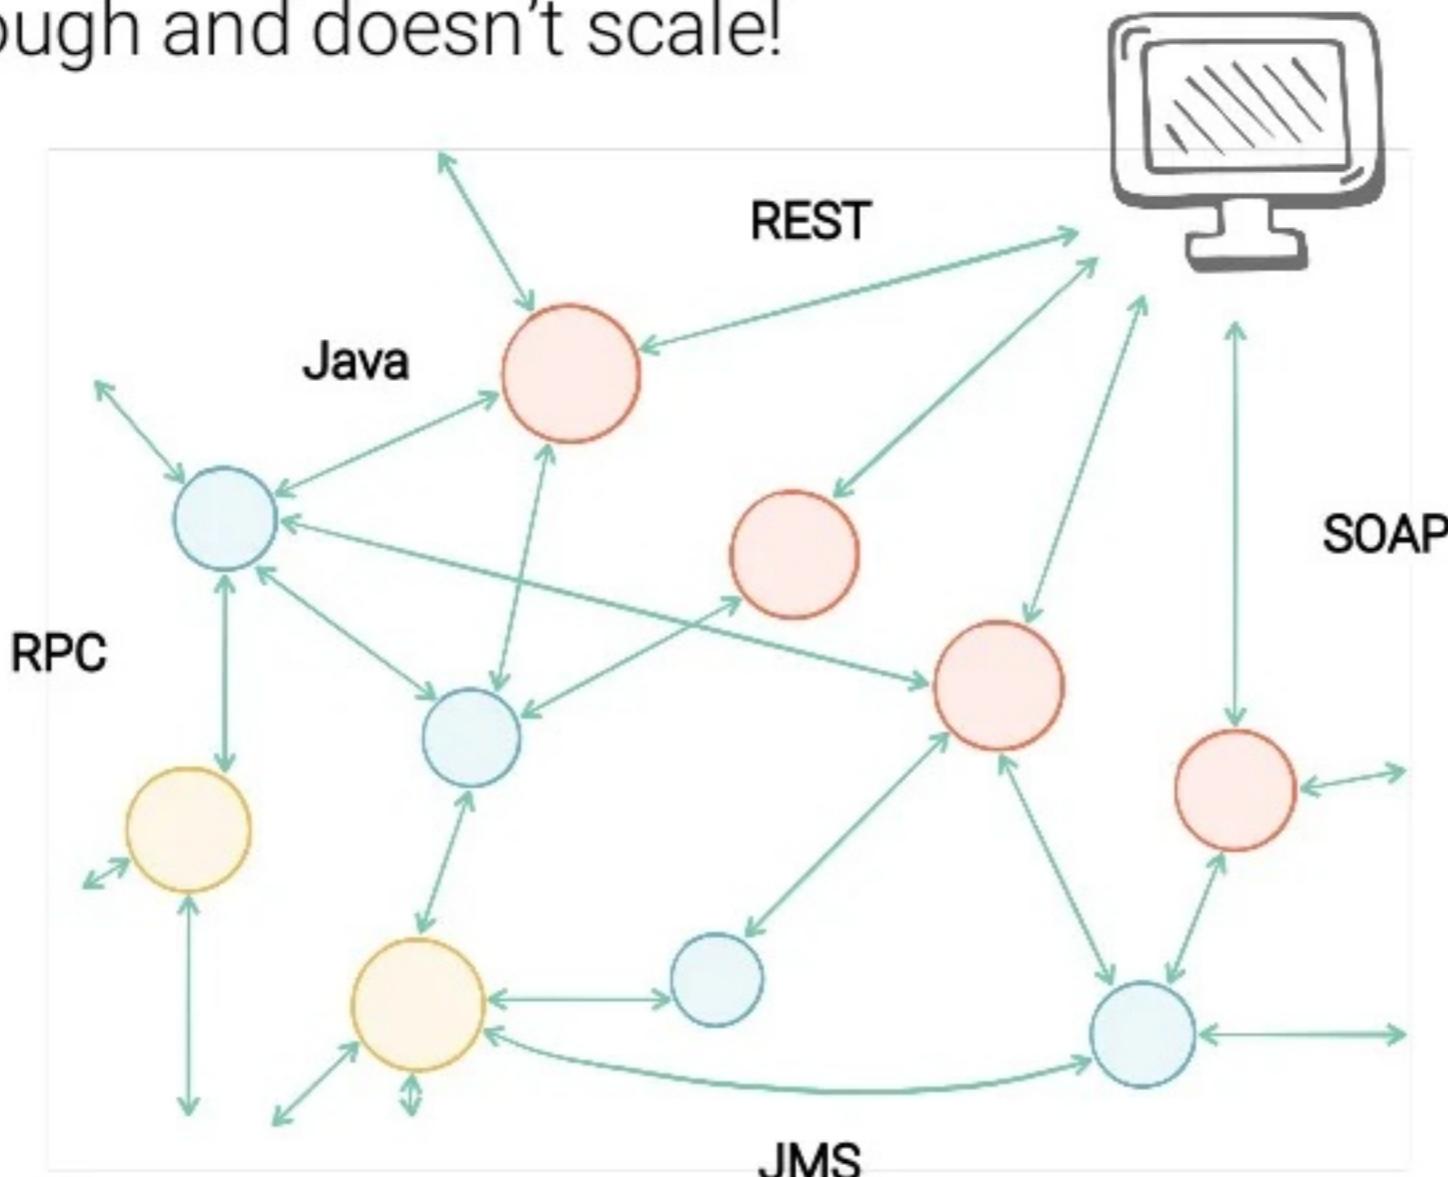
Can we ever get away from  
all legacy applications?

Probably not!



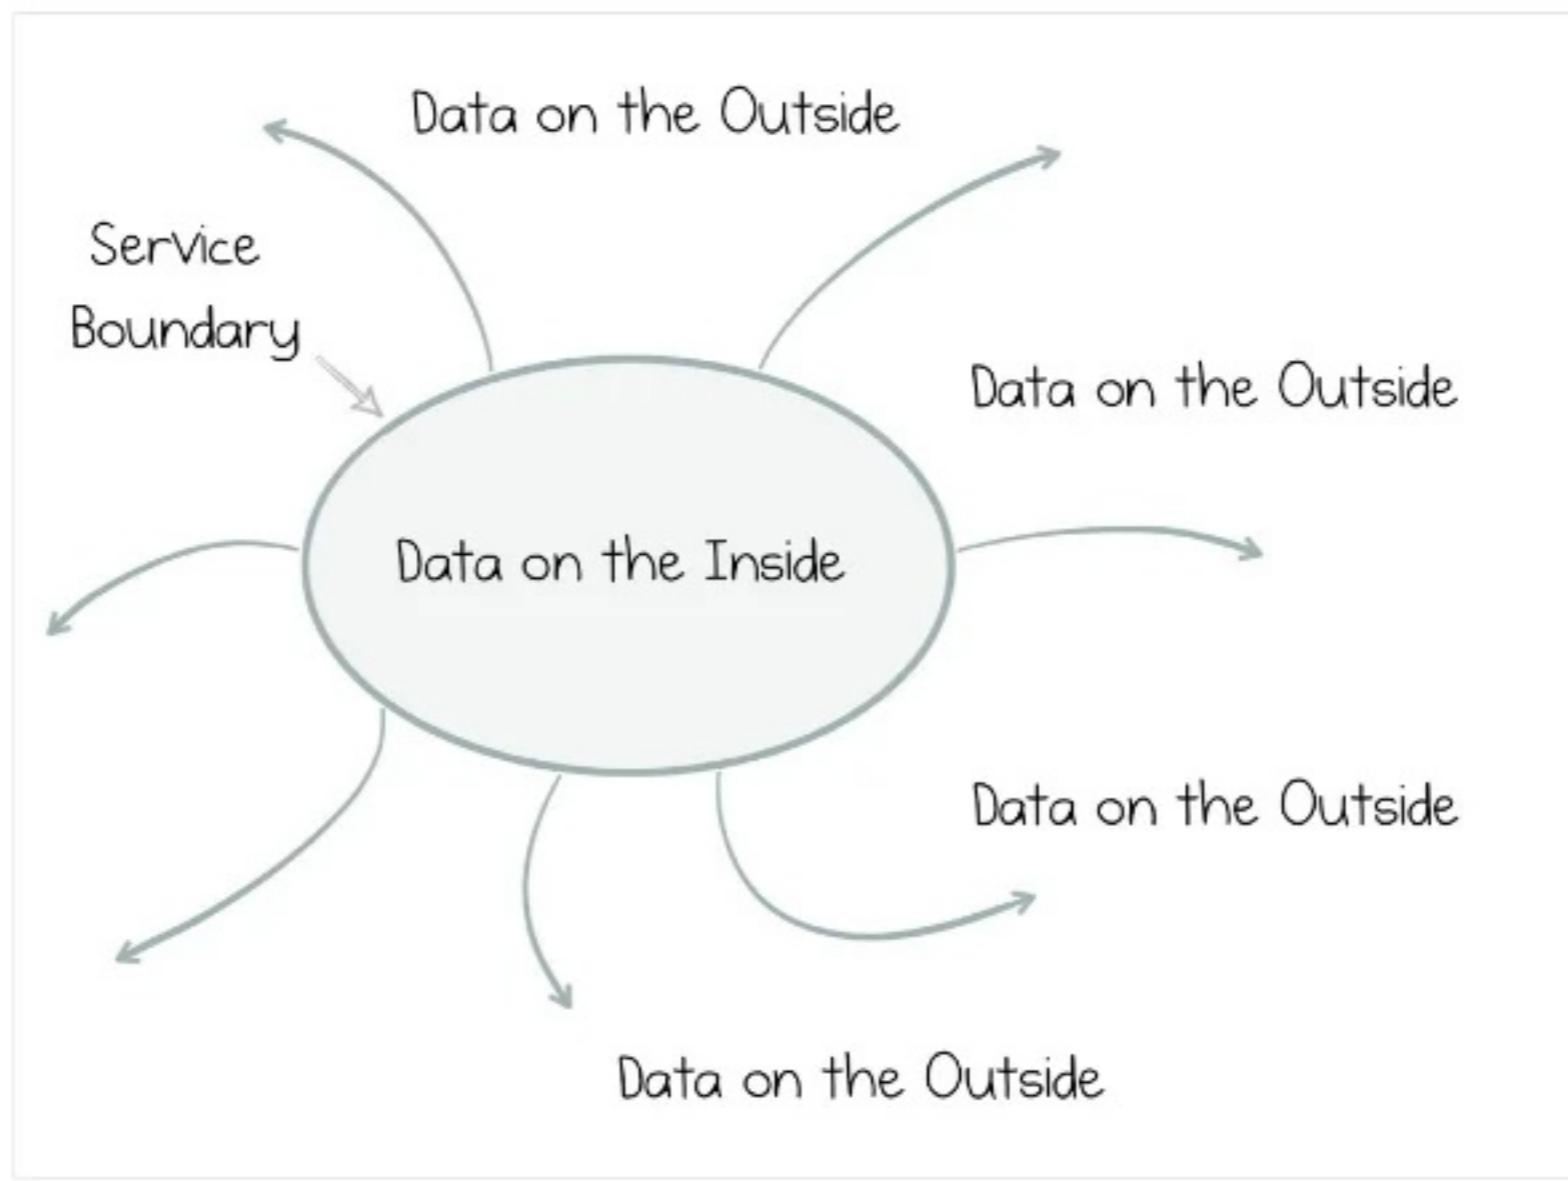
RPC / Request-Response... You can use the ESB for it. BUT:

... it often isn't enough and doesn't scale!

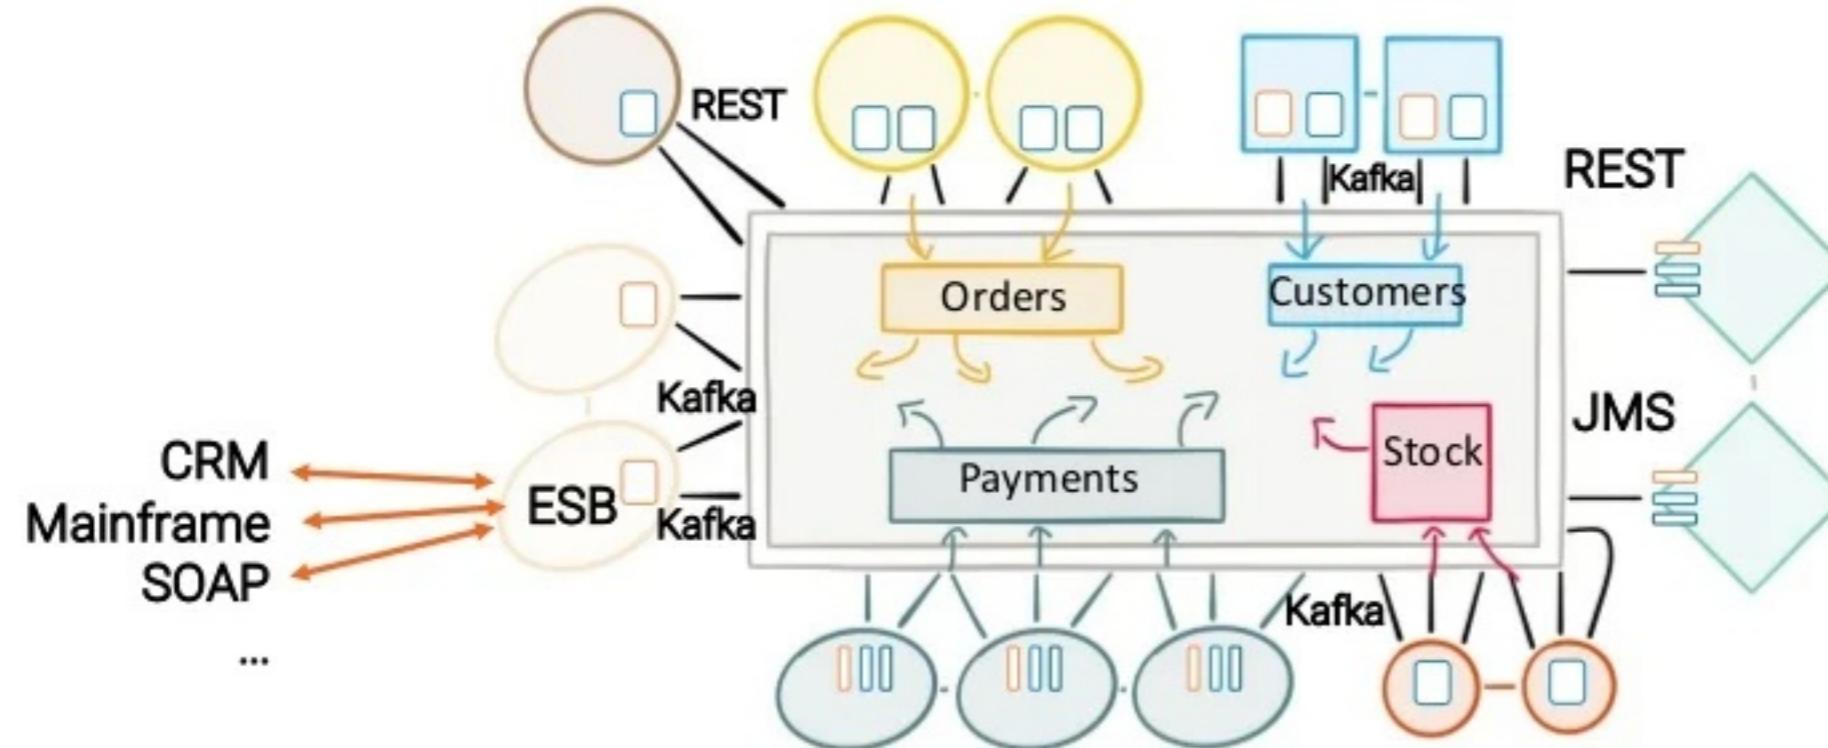


# Embrace data that lives and flows between services

---

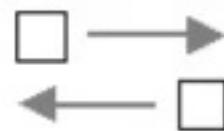


# An Event Streaming Platform gives services independence



Freedom to tap into and manage shared data

If you really need Request-Response with Kafka → A **Correlation ID** is your friend!

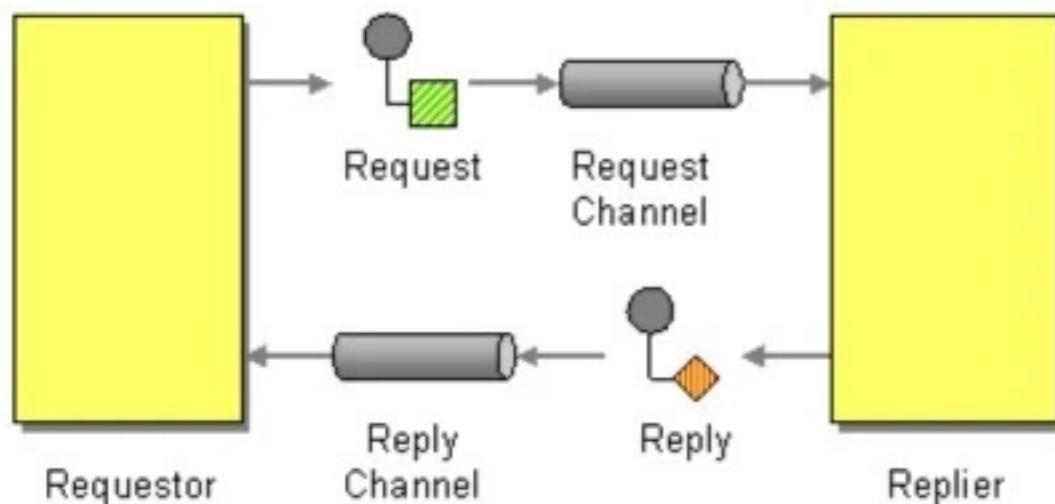


## Request-Reply

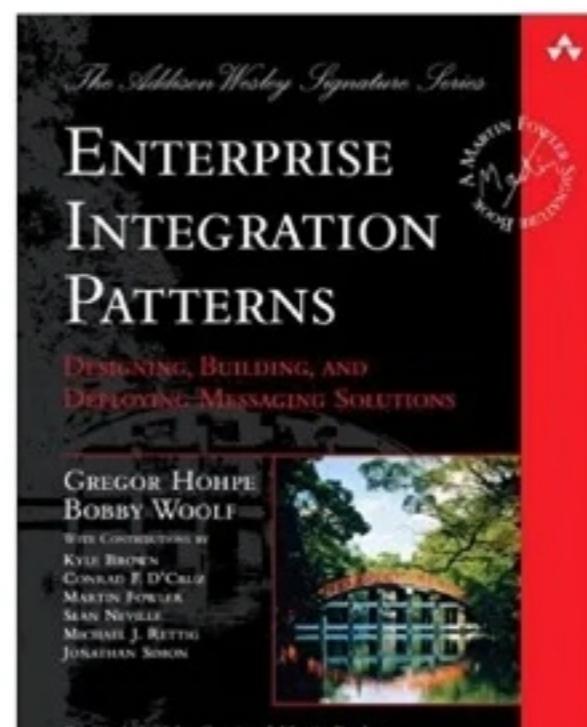
[MESSAGING PATTERNS](#) » [MESSAGE CONSTRUCTION](#) » REQUEST-REPLY

When two applications communicate via *Messaging*, the communication is one-way. The applications may want a two-way conversation.

**When an application sends a message, how can it get a response from the receiver?**

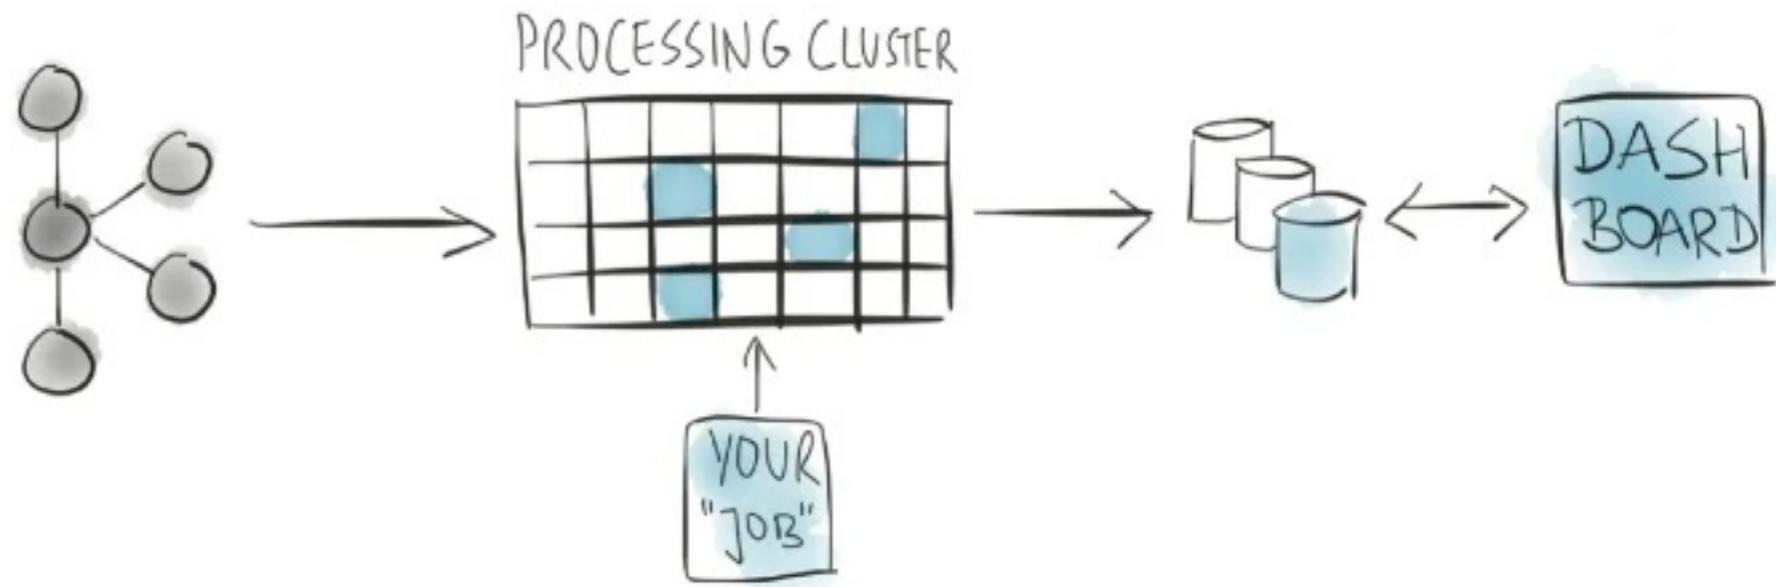


**Send a pair of *Request-Reply* messages, each on its own channel.**

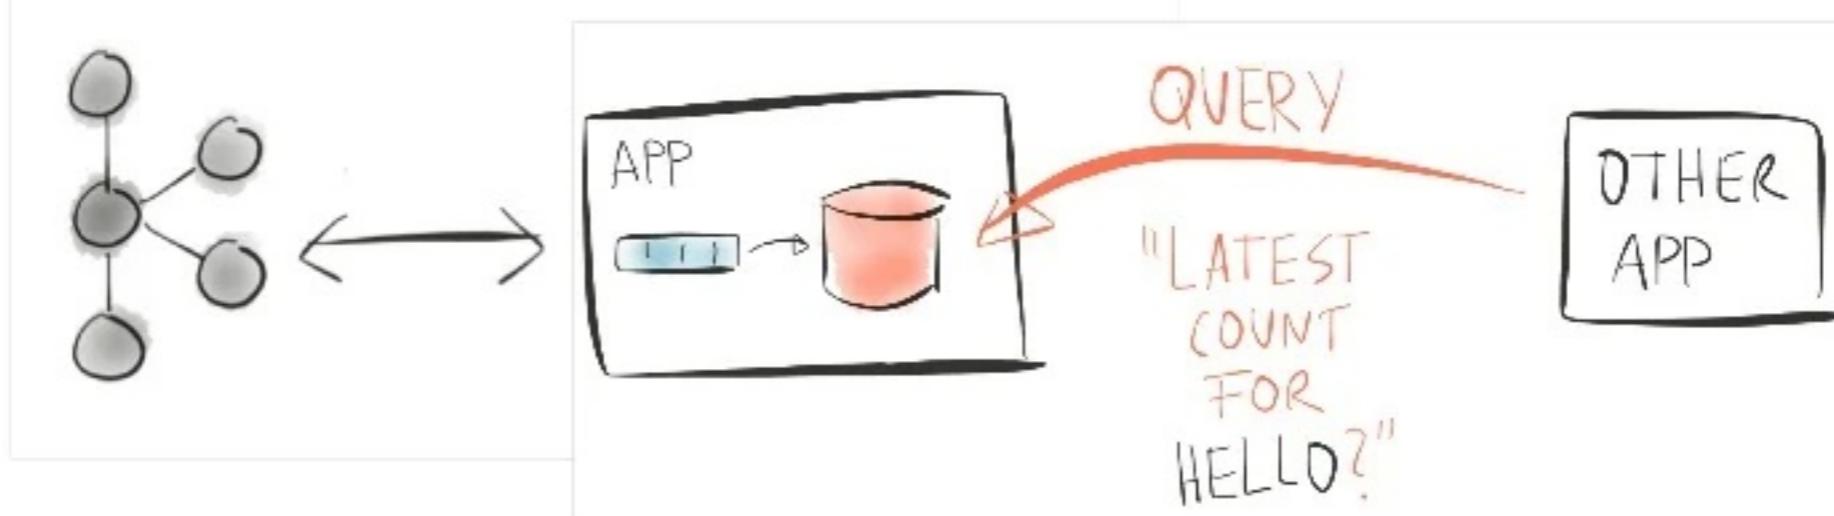


# If you need to combine RPC with Streaming: **Interactive Queries (IQ)** (Kafka Streams)

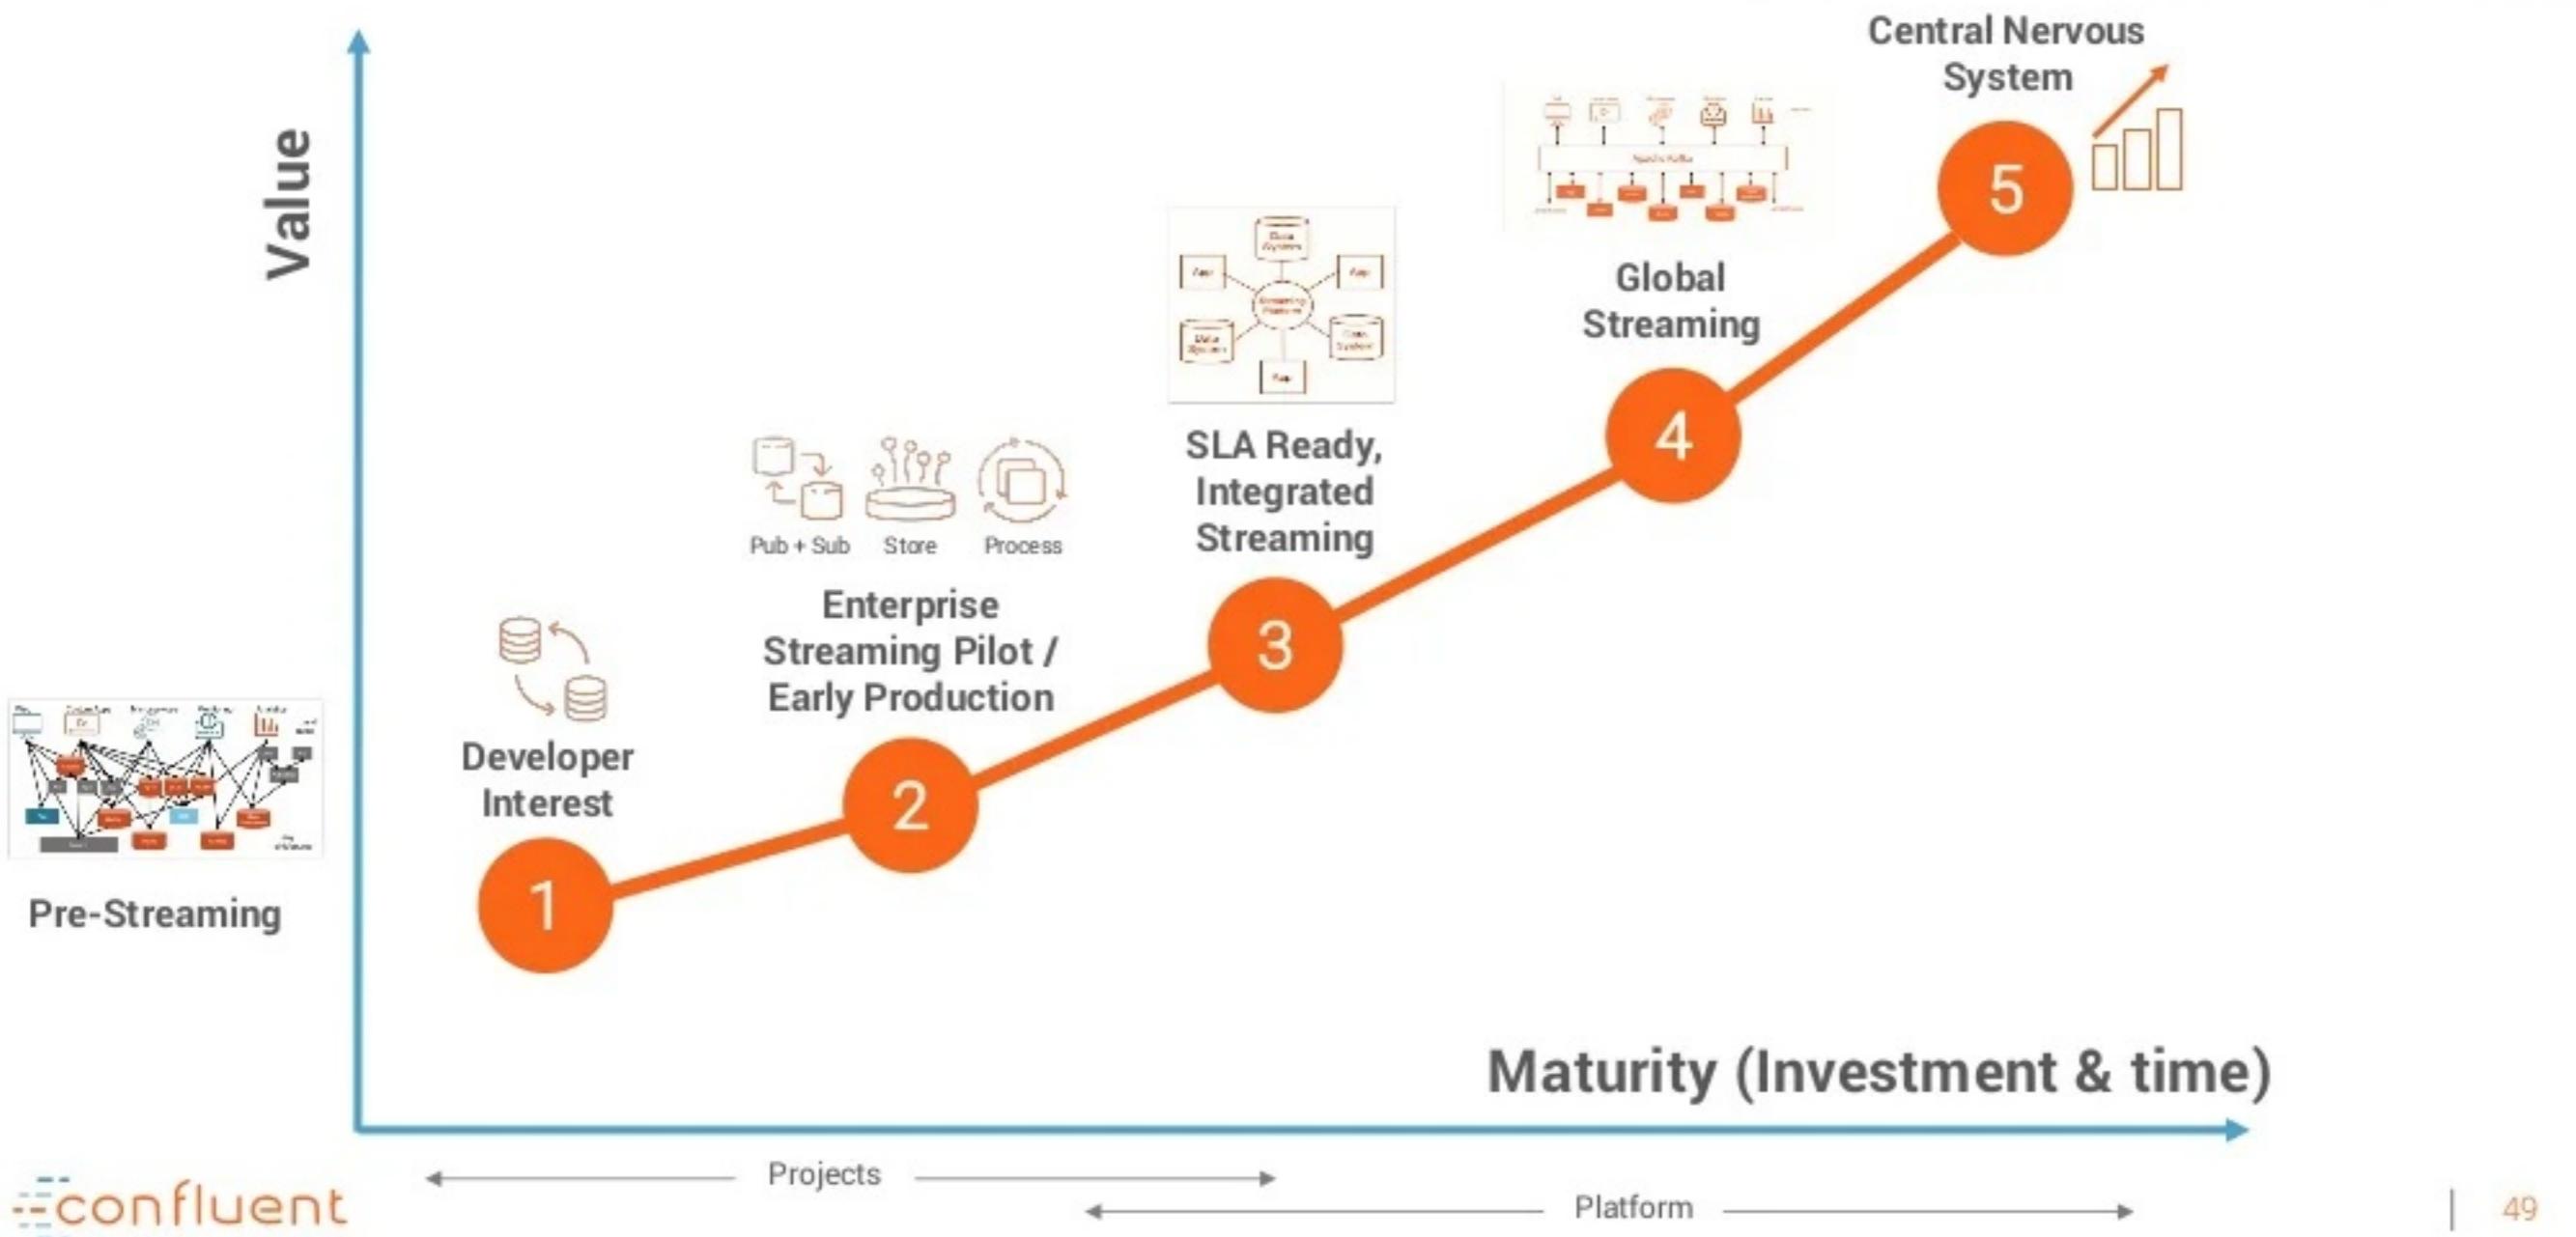
Before IQ ☹



With IQ ☺



# Confluent's Streaming Maturity Model - where are you?



# Frenemies?

---



How to integrate  
the old and new world?

# Mainframe offloading to Apache Kafka

Transaction Data

Date	Amount
1/27/2017	\$4.56
1/22/2017	\$32.14

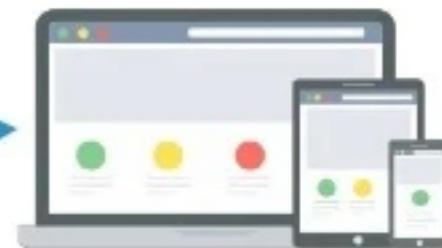
Transaction Description

Vendor	Description
Starbucks	Coffee
Walmart	Blu-Ray

Integration via

- Kafka Connect (JMS, MQ)
- REST Proxy
- 3<sup>rd</sup> party CDC tool
- Etc.

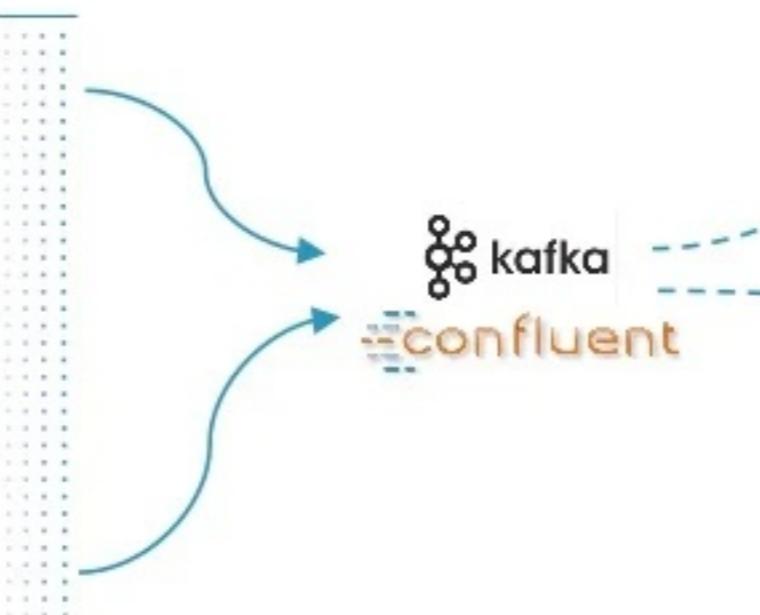
Website



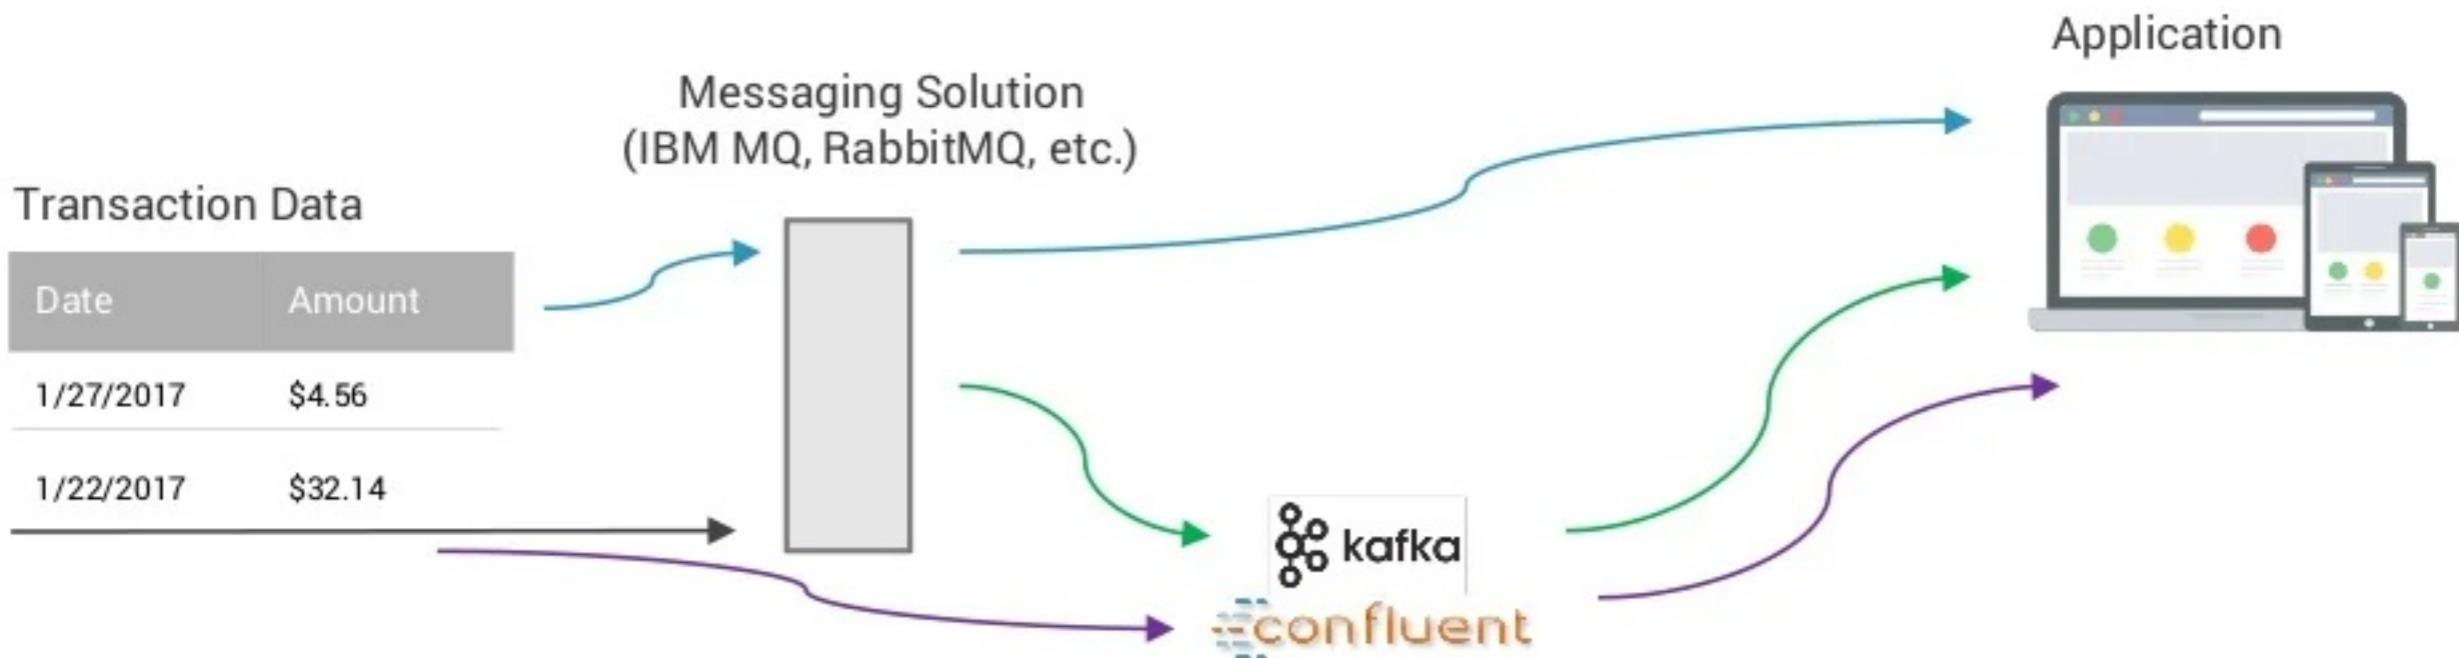
Client profiles



Mainframe MIPS = \$\$

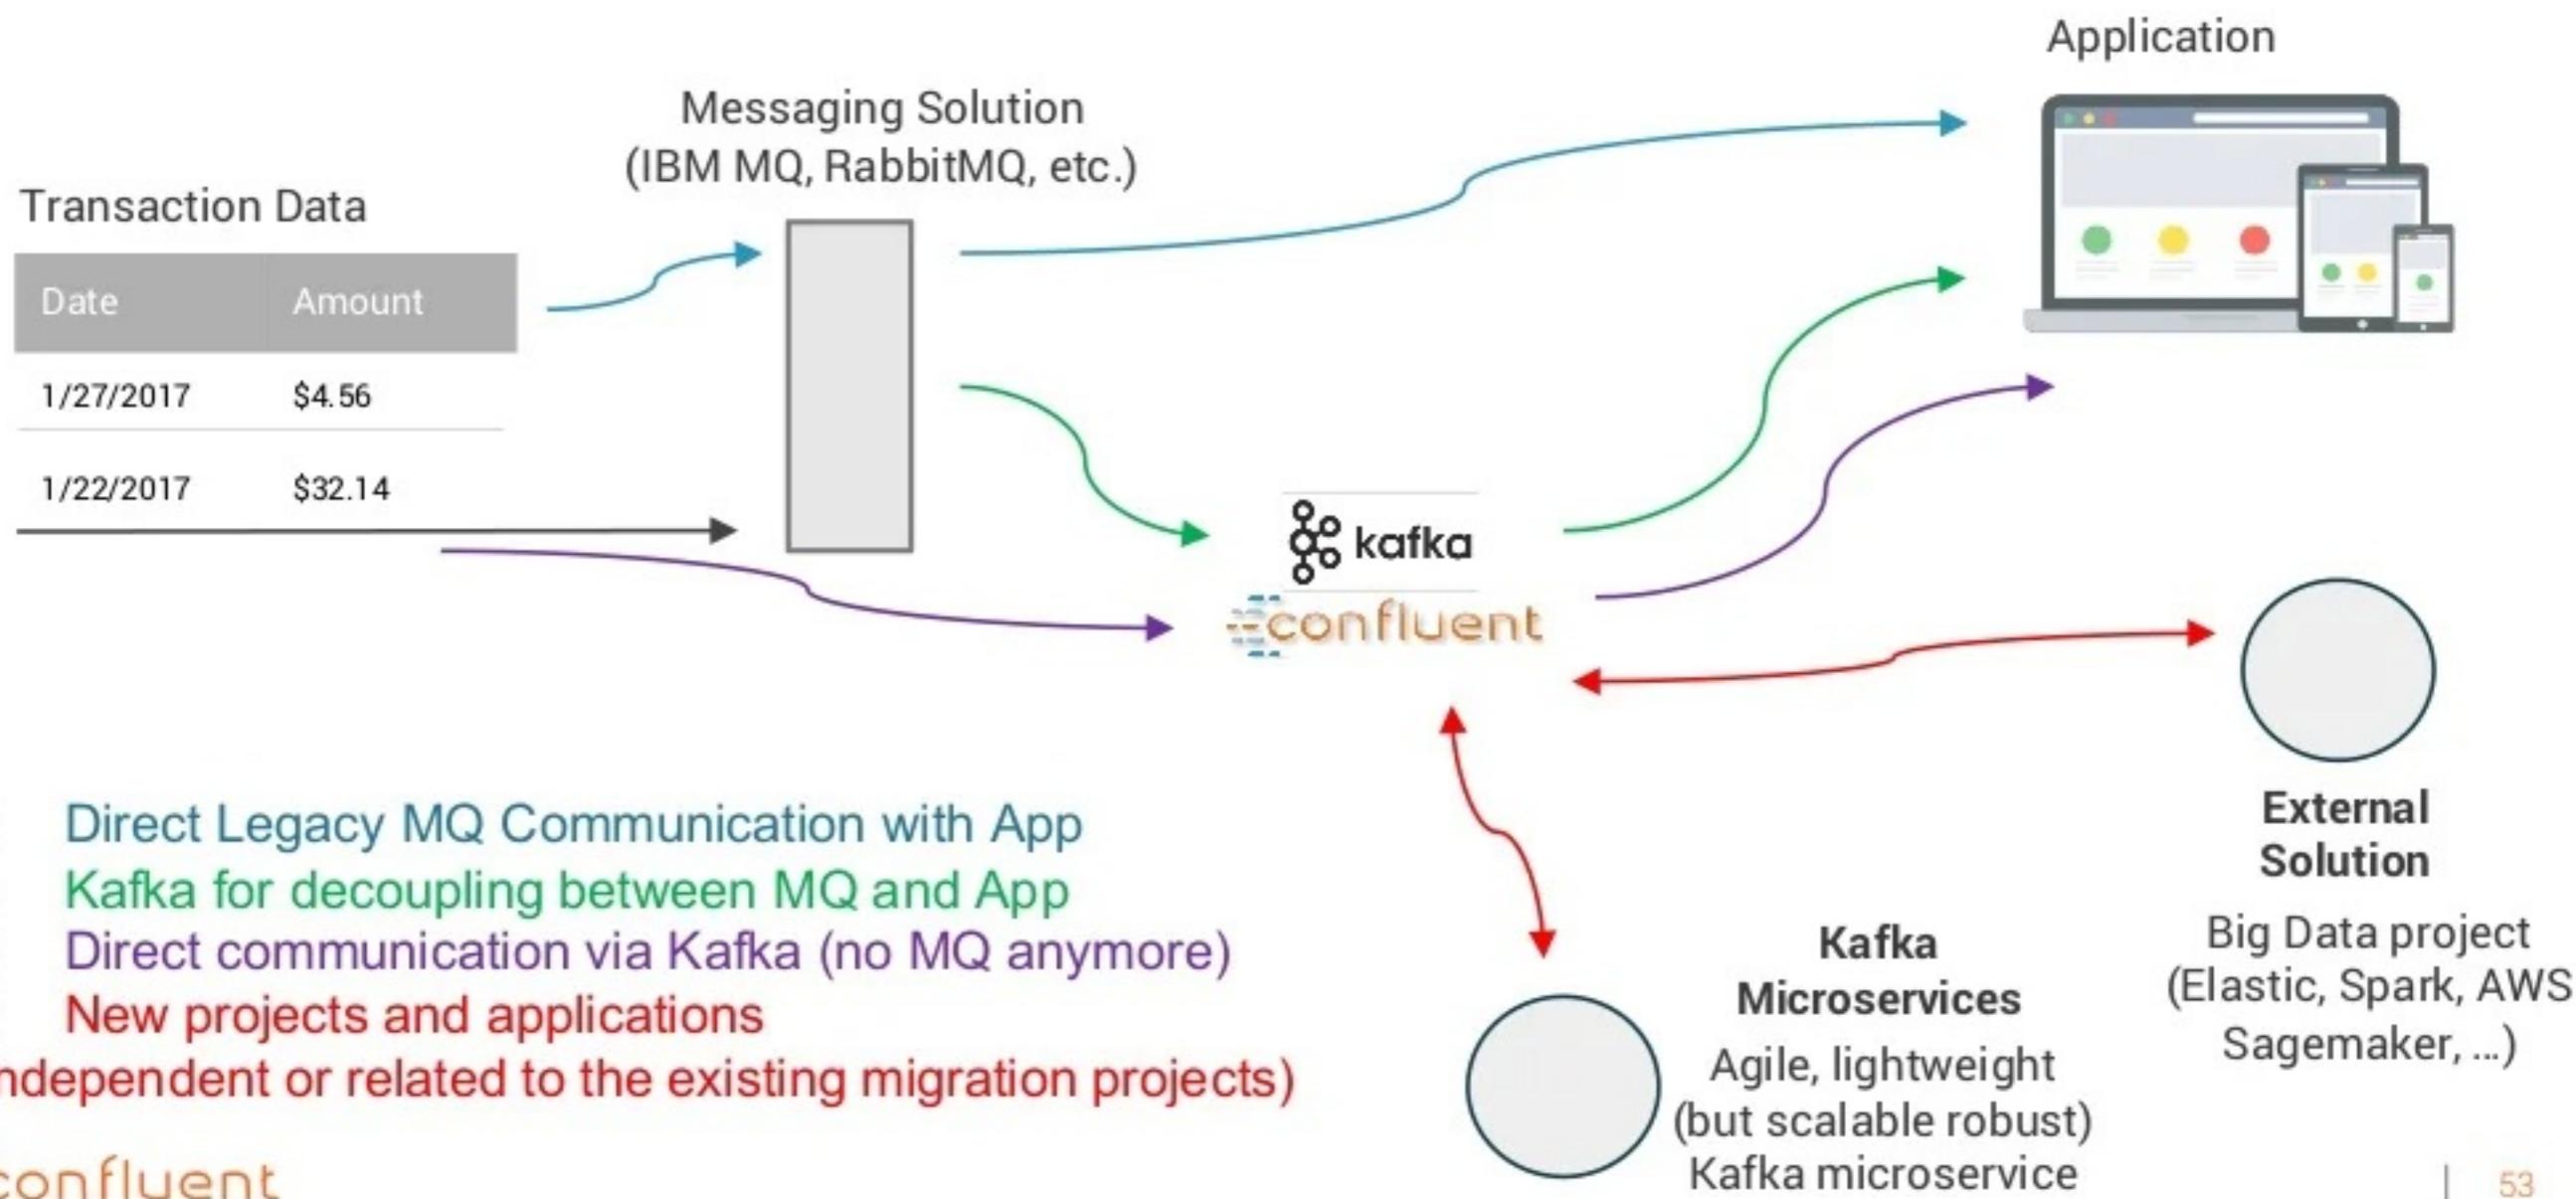


# MQ Integration (and later Replacement)



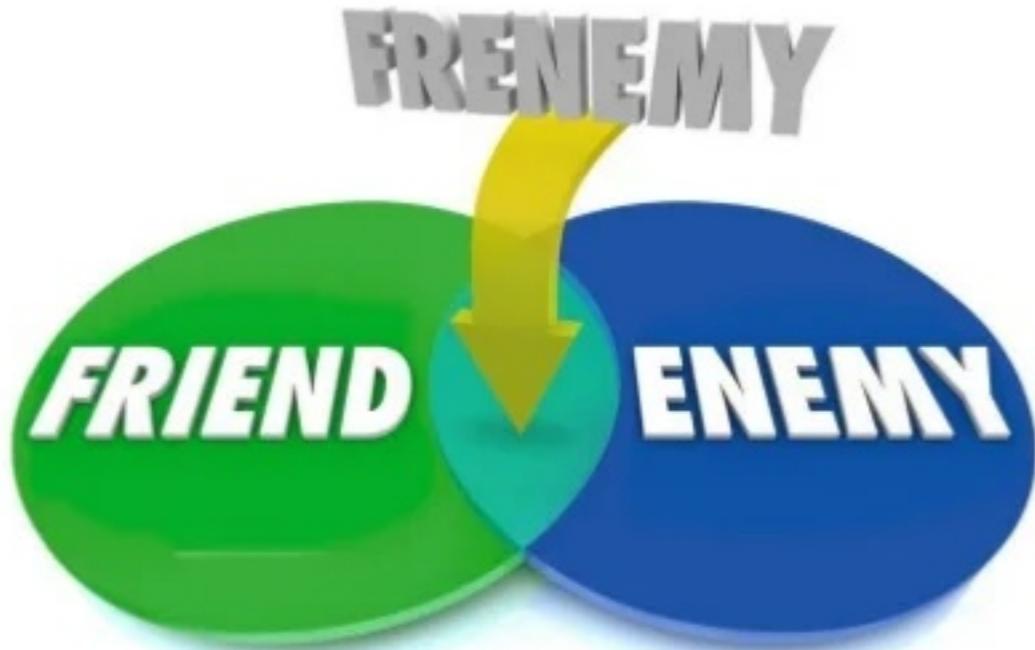
- 1) Legacy MQ Communication with App
- 2) Kafka for decoupling between MQ and App
- 3) Direct communication via Kafka (no MQ anymore)

# New, Innovative Projects and Applications



Use the right tool for the job (and combine them where it makes sense!)

---



This is just the beginning of a new era! Confluent Vision for Kafka:

---

### Global

Automated **disaster recovery**

Global applications with **geo-awareness**

### Infinite

Efficient and infinite data with **tiered storage**

Unlimited horizontal scalability for single clusters

Faster elastic scaling for brokers and partition

### Elastic

Easy Container-based orchestration and management

Faster **elastic scaling** when adding brokers and partitions

## Cloud-native Apache Kafka

# Questions? Feedback? Let's connect!

Kai Waehner

Technology Evangelist

[kontakt@kai-waehner.de](mailto:kontakt@kai-waehner.de)

[@KaiWaehner](https://twitter.com/KaiWaehner)

[www.confluent.io](http://www.confluent.io)

[www.kai-waehner.de](http://www.kai-waehner.de)

[LinkedIn](#)

